

**CMSC 409:  
Artificial Intelligence**

<http://www.people.vcu.edu/~mmanic/>

**Virginia Commonwealth University,  
Fall 2023,  
Dr. Milos Manic  
([mmanic@vcu.edu](mailto:mmanic@vcu.edu))**

1

**CMSC 409: Artificial Intelligence**

**Session # 21&22**

**Topics for today**

- Announcements
- Project 3, lessons learned
- Previous session review
- Error Back Propagation (EBP) algorithm
  - *Derivation*
  - *Learning example*

**Upcoming topics...**

- Probability Theory
- Subjective probability - Bayes' Rule

2

## CMSC 409: Artificial Intelligence

### Announcements Session # 21&22

- **IMPORTANT:**
  - Course materials (slides, assignments) are copyrighted by instructor & VCU. Sharing/posting/chatGPT/similar is copyright infringement and is strictly prohibited. Such must be immediately reported.
- Canvas
  - Prev. session slides updated
- TAs
  - Victor Cobilean <cobileanv@vcu.edu>, Harindra Sandun Mavikumbure mavikumbureh@vcu.edu
  - TA office hours: Thursdays, 3:30 - 4:30pm (Zoom)
- Project #3
  - Deadline was Oct. 26; Review a week from the deadline
- Project #4
  - Deadline is Nov. 9
- Paper (optional)
  - The 3<sup>rd</sup> draft due Nov. 3 (noon)
  - In addition to previous draft, it should contain a technique (or selection thereof), you plan on using to solve the selected problem (check out the class paper instructions for the 3<sup>rd</sup> draft)
  - The 4<sup>th</sup> draft (final submission) due Nov. 28
  - In addition to previous draft, it should contain a technique (or selection thereof), you plan on using to solve the selected problem (check out the class paper instructions for the 4<sup>th</sup> draft)
- Subject line and signature
  - Please use [CMSC 409] Last\_Name Question

3

## CMSC 409: Artificial Intelligence

### Session # 21&22

### Announcements, cont.

- Nov. 8, VCU closed
- **Final exam**
  - Dec. 13 (take home, 48 hr open book exam)
  - prep examples will be posted prior
- Interest in PhD or MSc program?
- Subject line and signature
  - Please use [CMSC 409] Last\_Name Question

4



**2022 Cybersecurity Speaker Series**

**Cyber security and future workforce in defending information technology infrastructure**

**November 4, 2022 11 a.m.–Noon (Eastern Time)**  
 Hosted by **Milos Manic, Ph.D., FIEEE, FCCI**  
 VCU College of Engineering West Hall, Room 106



**Major General Jan C. Norris**  
 Deputy Chief Information Officer (IMA)  
 The Office of the Chief Information Officer  
 United States Army Reserve

Register at:  
<https://cybersecurity.vcu.edu>

*Session 21&22, Updated on 11/1/23 11:52:56 AM*

5

# Lessons learned

## Project 03

© M. Manic, CMSC 409: Artificial Intelligence, F23

*Session 21&22, Updated on 11/1/23 11:52:56 AM*

6

## Class Statistics

### STATISTICS

COUNT	24
Minimum Value	0
Maximum Value	18
Range	18
Average	14.91
Median	17
Standard Deviation	5.03
Variance	25.38

### GRADE DISTRIBUTION

Greater than 100	13
90 - 100	6
80 - 89	3
70 - 79	0
60 - 69	0
50 - 59	0
40 - 49	0
30 - 39	0
20 - 29	0
10 - 19	0
0 - 9	2

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:56 AM

7

## Project 3 review

### Pr. 3.1 Understand and explore Reinforcement Learning (6pts)

1. Discuss what Reinforcement Learning (RL) is. (2pts)
2. What are the real-world applications which use RL, discuss 2 application areas. (2pts)
3. Discuss the advantages and disadvantages of RL. (2pts)

### Best practices:

#### Provide complete answers:

- Concise and to the point answers
- Try to avoid lengthy and possibly unrelated discussions

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:56 AM

8

## Project 3 review

### Pr. 3.2 Simulation of a self-driving cab using Q Learning (9 pts)

1. Change the following hyperparameters of the RL pipeline. Make a note of the reward for each combination of the hyperparameters. Discuss how the values of the hyperparameters affect the total reward for the smart cab **(6pts)**

- a. alpha: 0.01, 0.1, 0.9
- b. gamma: 0.40, 0.99
- c. epsilon: 0.5, 1

**Example:** You can try:

- alpha: 0.01
- gamma: 0.40
- epsilon: 0.5

(then change alpha for the same gamma and epsilon, etc.)

2. Discuss the importance of reducing the epsilon parameter in achieving a higher reward **(1pts)**
3. Run the “*Smart\_cab\_brute\_force.ipynb*” and compare and discuss the results with the reinforcement learning results. **(2pts)**

## Project 3 review

### Pr.3.2

#### Importance of parameters

$$Q_{t+1}(s+a) = Q_t(s,a) + \alpha(R + \lambda \max Q_t(s+1, :) - Q_t(s,a))$$

#### $\alpha$ - Learning Rate:

- How much  $R + \lambda \max Q_t(s+1, :) - Q_t(s,a)$  contributes to *the*  $Q$ -values in the next time step,  $Q_{t+1}(s+a)$ :
  - Define the portion of reward  $R$ , discount factor  $\lambda$ , prev. step  $Q$ -value  $Q_t(s,a)$ , and  $\max Q$
  - $\max Q$  or  $\max Q_t(s+1, :)$ , evaluates all possible actions for “neighboring” state  $s+1$ .
- Larger alpha  $\rightarrow$  faster convergence,
  - but may lead to convergence challenges (the “convergence” here refers to the process of algorithm converging to the maximum reward -  $R$ )
  - As algorithm goes through iterations, it attempts to maximize reward for given  $Q(\text{state}, \text{action})$ ,
  - But, with larger  $\alpha$ , algorithm may get stuck with lower reward (may not be able to improve reward). But, the process also depends on  $\lambda$  and epsilon...

## Project 3 review

### Pr.3.2

#### Importance of parameters

- $\lambda [0, 1]$  - Discount Factor
  - larger lambda (closer to 1) - makes agent focus on the long-term effective award;
  - smaller lambda (closer to 0) - makes agent focus on immediate reward (greedy agent).

## Project 3 review

### Pr.3.2

#### Importance of parameters

- $\mathcal{E}_{decay}[0, 1]$  - decay rate
  - *Epsilon closer to 1* - agent performs more explorations (more random actions)
  - *Epsilon closer to 0* - agent performs more exploitation (follows  $Q$ -table).

#### Exploration vs. exploitation

- With exploration, an agent takes random decisions; Useful when learning about the environment.
- With exploitation, an agent takes actions based on what agent already knows (from the  $Q$  table)

## Project 3 review

Alpha	Gamma	Epsilon	Reward
0.01	0.99	0.5	6.09
0.01	0.99	1	6.8
0.01	0.4	0.5	0.5
0.01	0.4	1	2.4
0.1	0.99	0.5	8.5
0.1	0.99	1	8.7
0.1	0.4	0.5	7.65
0.1	0.4	1	8.2
0.9	0.99	0.5	8.2
0.9	0.99	1	6.9
0.9	0.4	0.5	7.55
0.9	0.4	1	7.3

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:57 AM

13

## Project 3 review

### Extra Credit: (3pts)

Implement the *epsilon\_decay* as:

$$\epsilon_{decay} = \min\_epsilon + (\max\_epsilon - \min\_epsilon) * \exp(-\epsilon_{decay\_rate} * episode)$$

using the following hyperparameters:

- min\_epsilon
- max\_epsilon
- epsilon\_decay\_rate

1. Which values of these hyperparameters (alpha, gamma, epsilon, min\_epsilon, max\_epsilon, epsilon\_decay\_rate) result in the highest reward?
2. What is the highest reward?

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:57 AM

14

## Project 3 review

- **Extra Credit**

```
# Implementing the Q Learning Algorithm:
for episode in range(total_ep):
    # Reset Environment:
    state = env.reset()
    step = 0
    done = False

    for step in range(max_steps):
        # Choose an action a in the current world state(s) (step 3)
        # First we randomize a number
        exp_exp_tradeoff = random.uniform(0, 1)

        # If this number > greater than epsilon --> exploitation (taking the biggest q value for the current state):
        if exp_exp_tradeoff > epsilon:
            action = np.argmax(q_table[state, :])

        # Else, doing random choice:
        else:
            action = env.action_space.sample()

        # Take the action (a) and observe the outcome state (s') and the reward (r)
        new_state, reward, done, info = env.step(action)

        # Update Q(s,a) = Q(s,a) + lr [R(s,a) + gamma * max Q(s',a') - Q(s,a)]
        q_table[state, action] = q_table[state, action] + lr * (reward + gamma *
            np.max(q_table[new_state, :]) - q_table[state, action])

        # Our new state:
        state = new_state

        # If done True, finish the episode:
        if done == True:
            break

    # Reduce epsilon (because we need less and less exploration):
    epsilon = min_epsilon + (max_epsilon - min_epsilon) * np.exp(-decay_rate*episode)
    # Increment number of episodes:
    episode += 1

print("Training finished.\n")
```

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:57 AM

### Exploitation vs. Exploration

- As agent learns, epsilon parameter gets smaller
- Through the process, an agent gradually moves from exploration to exploitation

15

## During submission

- Project deliverable should be a zip file containing:
  - Written report (pdf) with answers to all of the questions above.
  - Updated Source code in Ipython notebook file (.ipynb)
- Submit your zip file to Canvas. Please name the zip file as GroupName\_Project3.zip.

© M. Manic, CMSC 409: Artificial Intelligence, F23

Session 21&22, Updated on 11/1/23 11:52:57 AM

16



## Error Back Propagation (EBP) algorithm

### ➡ *Derivation*

- *Learning example*
- *Heuristic approaches to Error Back Propagation modifications*
  - *variable gain, alpha, weight rescaling,*
  - *momentum,*
  - *search along the gradient*
  - *Quickprop, RPROP, Delta-Bar-Delta, Back Percolation*

## Error-Back Propagation algorithm

- Werbos 1974,
- Later Rumelhart & McClelland 1986.

### Error-Backpropagation Algorithm (EBP)

- *Breakthrough, finally annulling Minsky & Papert's influence from their book "Perceptrons" 1969.*
- *Problems:*
  - *Convergence difficulties*
  - *Oscillations*
  - *Speed*


[Werbos 74] Paul J. Werbos. *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences*. PhD thesis, Harvard University, 1974.

[Werbos 94] Paul John Werbos, "The Roots of Backpropagation: From Ordered Derivatives to Neural Networks and Political Forecasting, Wiley-Interscience, January 1994.

[Rumelhart 86] Rumelhart, D.E., McClelland, J.L., *Parallel Distributed Processing: Explorations in the Microstructure of Cognition, Vol.1*, Cambridge, Ma, MIT Press, 1986.

## EBP Error Back Propagation algorithm

single output



where:

$f(x)$  – activation function

$z_p$  – neuron outputs

$x_p$  – input pattern

$F\{z\}$  – non-linear function based on activation function of several neurons


For each pattern  $p$ ,  
network output is:

$$o_p = F\left\{f\left(w_1x_{p1} + w_2x_{p2} + \cdots + w_nx_n\right)\right\}$$

$$\text{Total error: } TE = \sum_{p=1}^{np} [d_p - o_p]^2$$

19

## EBP Error Back Propagation algorithm



single output case...

$$TE = \sum_{p=1}^{np} [d_p - o_p]^2$$

$$o_p = F\left\{f\left(w_1x_{p1} + w_2x_{p2} + \cdots + w_nx_{pn}\right)\right\}$$

Consider a single weight  $w_i$ . The gradient of  $TE$  along  $w_i$  is:

$$\frac{d(TE)}{dw_i} = -\sum_{p=1}^{np} \left[ 2(d_p - o_p) \frac{do_p}{dz_p} \frac{dz_p}{dnet_p} \frac{dnet_p}{dw_i} \right]$$

(goal is to minimize  $TE$ )

20

## EBP Error Back Propagation algorithm

The gradient of error function with respect to weights:

$$\frac{d(TE)}{dw_i} = - \sum_{p=1}^{np} \left[ 2(d_p - o_p) \frac{do_p}{dz_p} \frac{dz_p}{dnet_p} \frac{dnet_p}{dw_i} \right]$$

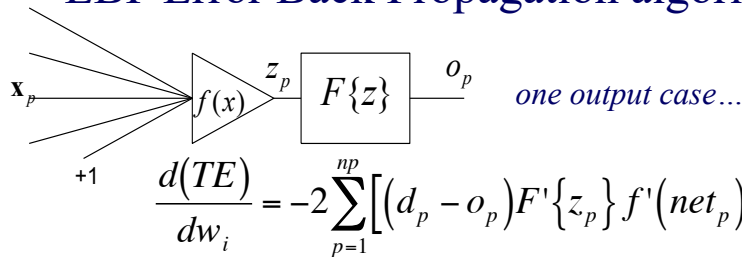
$$\frac{d(TE)}{dw_i} = -2 \sum_{p=1}^{np} \left[ (d_p - o_p) F'\{z_p\} \cdot f'(net_p) \cdot x_{pi} \right]$$

$\frac{do_p}{dz_p}$  - gain of the function  $F$  through the network       $\frac{dnet_p}{dw_i}$  - change of net with regards to weight

$\frac{dz_p}{dnet_p}$  - slope of activation function

21

## EBP Error Back Propagation algorithm



To move towards the minimum, the weight change should be opposite to the gradient change (direction).

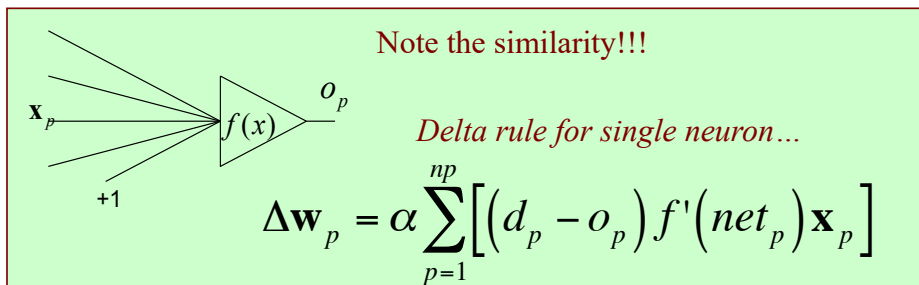
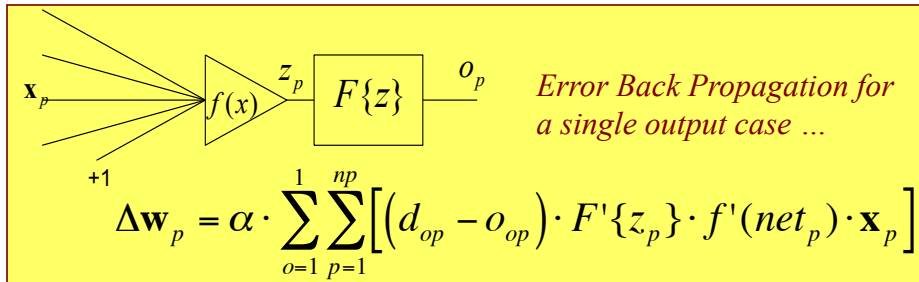
$$\Delta w_{pi} \sim 2 \sum_{p=1}^{np} \left[ (d_p - o_p) \cdot F'\{z_p\} \cdot f'(net_p) \cdot x_{pi} \right]$$

...which can be extended for all weights of this neuron ( $\alpha$  instead of ~):

$$\Delta \mathbf{w}_p = \alpha \cdot \sum_{o=1}^{no} \sum_{p=1}^{np} \left[ (d_{op} - o_{op}) \cdot F'\{z_p\} \cdot f'(net_p) \cdot \mathbf{x}_p \right]$$

22

## EBP Error Back Propagation algorithm




© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 23

Session 21&22, Updated on 11/1/23 11:52:57 AM

23

## EBP Error Back Propagation algorithm



The weight increment for the whole network is a **superposition** of weight modifications from errors on all outputs:

$$\Delta \mathbf{w}_p = \alpha \cdot \sum_{o=1}^{no} \sum_{p=1}^{np} \left[ (d_{op} - o_{op}) \cdot F'\{z_p\} \cdot f'(net_p) \cdot \mathbf{x}_p \right]$$

*...weight update by each error from each output....*

© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 24

Session 21&22, Updated on 11/1/23 11:52:57 AM

24

## EBP Error Back Propagation algorithm

$$\Delta \mathbf{w}_p = \alpha \cdot \sum_{o=1}^{no} \sum_{p=1}^{np} [(d_{op} - o_{op}) \cdot F'\{z_p\} \cdot f'(net_p) \cdot \mathbf{x}_p]$$

**A layer**

**B layer**

$\Delta \mathbf{w} = \alpha \cdot \Delta \cdot \mathbf{x}$

- Example of 2-layer network
- Forward computation & backward computation
- Error is computed back (back propagated)

© M. Manic, CMSC 409: Artificial Intelligence, F23      Page 25      Session 21&22, Updated on 11/1/23 11:52:57 AM

25

## Error Back Propagation (EBP) algorithm

- Derivation
- ➡ Learning example
- Heuristic approaches to Error Back Propagation modifications
  - variable gain, alpha, weight rescaling,
  - momentum,
  - search along the gradient
  - Quickprop, RPROP, Delta-Bar-Delta, Back Percolation


© M. Manic, CMSC 409: Artificial Intelligence, F23      Page 26      Session 21&22, Updated on 11/1/23 11:52:57 AM

26

## EBP Error Back Propagation algorithm

**EXAMPLE:**

$$k = 0.5 \quad \alpha = 1.0$$



• Example of 2-layer network

• We will process single pattern

• with parameters chosen as:

$$\mathbf{x} = \begin{bmatrix} +1 \\ -1 \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} +1 & +2 & +.5 \\ -1 & +1 & -.5 \\ +1 & -1 & +2 \\ +2 & -2 & +1 \end{bmatrix} \quad \mathbf{d} = \begin{bmatrix} +1 \\ -1 \end{bmatrix}$$

27

## EBP Error Back Propagation algorithm


*example*

**1. Forward Computation**

$$k = 0.5 \quad \alpha = 1.0$$

$$\text{net} = \sum w_i x_i$$

$$\text{out} = \tanh(k \cdot \text{net})$$



$$\text{net}_1 = 1 \cdot 1 - 1 \cdot 2 + 1 \cdot 0.5 = -0.5;$$

$$\text{out}_1 = \tanh(0.5 \cdot \text{net}_1) = -0.2449$$

$$\text{net}_2 = -2.5;$$

$$\text{out}_2 = \tanh(0.5 \cdot \text{net}_2) = -0.8483$$

now propagate these values to a 2<sup>nd</sup> layer.....

$$\text{net}_3 = 1 \cdot (-0.2449) - 1 \cdot (-0.8483) + 2 \cdot 1 = 2.6034;$$

$$\text{out}_3 = 0.86215$$

$$\text{out}_4 = 0.8017$$

28

## EBP Error Back Propagation algorithm example


### 2. Error Computation

$$d = \begin{bmatrix} +1 \\ -1 \end{bmatrix}$$

*error = desired - actual output*

$$err_3 = 1 - 0.86215 = 0.1378$$

$$err_4 = -1 - 0.8017 = -1.8017$$



29

## EBP Error Back Propagation algorithm example

### 3. Error back propagation through last layer

*soft activation function:*

$$f' = \text{gain} \cdot (1 - \text{out}^2)$$


$$\Delta_B = f' \cdot \text{err}$$

$$f'_3 = 0.5 \cdot (1 - 0.86215^2) = 0.1283$$

$$f'_4 = 0.5 \cdot (1 - 0.8017^2) = 0.1786$$

$$\Delta_3 = 0.1283 \cdot 0.1378 = 0.0177$$

$$\Delta_4 = 0.1786 \cdot (-1.8017) = -0.3218$$



30

## EBP Error Back Propagation algorithm example

### 4. Error back propagation through the first layer

$$f' = \text{gain} \cdot (1 - \text{out}^2)$$


$$\Delta_A = f' \cdot \sum (w \cdot \Delta_B)$$

$$f'_1 = 0.5 \cdot (1 - 0.2449^2) = 0.4700$$

$$f'_2 = 0.5 \cdot (1 - 0.8483^2) = 0.1402$$

$$\Delta_1 = 0.4700 \cdot (1 \cdot 0.0177 + 2 \cdot (-0.3218)) = -0.2942$$

$$\Delta_2 = 0.1402 \cdot ((-1) \cdot 0.0177 + (-2) \cdot (-0.3218)) = 0.0878$$



© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 31

Session 21&22, Updated on 11/1/23 11:52:58 AM

31


## EBP Error Back Propagation algorithm example

### 5. Calculate layer A weight update

$$\Delta \mathbf{w} = \text{alpha} \cdot \Delta \cdot \mathbf{x}$$

$$\Delta_1 = -0.2942 \begin{cases} \Delta w_{11} = 1.0 \cdot (-0.2942) \cdot 1 = -0.2942 \\ \Delta w_{12} = 1.0 \cdot (-0.2942) \cdot (-1) = 0.294 \text{ Neuron 1} \\ \Delta w_{13} = 1.0 \cdot (-0.2942) \cdot 1 = -0.294 \end{cases}$$

$$\Delta_2 = 0.0878 \begin{cases} \Delta w_{21} = 1.0 \cdot 0.0878 \cdot 1 = 0.0878 \\ \Delta w_{22} = 1.0 \cdot 0.0878 \cdot (-1) = -0.0878 \text{ Neuron 2} \\ \Delta w_{23} = 1.0 \cdot 0.0878 \cdot 1 = 0.0878 \end{cases}$$



© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 32

Session 21&22, Updated on 11/1/23 11:52:58 AM

32



## EBP Error Back Propagation algorithm

### example

#### 6. Perform layer A weight update


$$\mathbf{W} = \mathbf{W} + \Delta \mathbf{W}$$

Neuron 1

$$\begin{cases} w_{11} = w_{11} + \Delta w_{11} = 1 - 0.2942 = 0.7058 \\ w_{12} = w_{12} + \Delta w_{12} = 2 + 0.2942 = 2.2942 \\ w_{13} = w_{13} + \Delta w_{13} = 0.5 - 0.2942 = 0.2058 \end{cases}$$

Neuron 2

$$\begin{cases} w_{21} = w_{21} + \Delta w_{21} = -1 + 0.0878 = -0.9122 \\ w_{22} = w_{22} + \Delta w_{22} = 1 - 0.0878 = 0.9122 \\ w_{23} = w_{23} + \Delta w_{23} = -0.5 + 0.0878 = -0.4122 \end{cases}$$



© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 33

Session 21&22, Updated on 11/1/23 11:52:58 AM

33

## EBP Error Back Propagation algorithm

### example

#### 7. Calculate layer B weight update


$$\Delta \mathbf{W} = \text{alpha} \cdot \Delta \cdot \mathbf{x}$$

Neuron 3

$$\Delta_3 = 0.0177 \begin{cases} \Delta w_{31} = 1.0 \cdot 0.0177 \cdot (-0.2449) = -0.0043 \\ \Delta w_{32} = 1.0 \cdot 0.0177 \cdot (-0.8483) = -0.0150 \\ \Delta w_{33} = 1.0 \cdot 0.0177 \cdot 1 = 0.0177 \end{cases}$$

Neuron 4

$$\Delta_4 = -0.3218 \begin{cases} \Delta w_{41} = 1.0 \cdot (-0.3218) \cdot (-0.2449) = 0.0788 \\ \Delta w_{42} = 1.0 \cdot (-0.3218) \cdot (-0.8483) = 0.2730 \\ \Delta w_{43} = 1.0 \cdot (-0.3218) \cdot 1 = -0.3218 \end{cases}$$



© M. Manic, CMSC 409: Artificial Intelligence, F23

Page 34

Session 21&22, Updated on 11/1/23 11:52:58 AM

34


## EBP Error Back Propagation algorithm

### example

#### 8. Perform layer B weight update

$$\mathbf{W} = \mathbf{W} + \Delta \mathbf{W}$$

$$\begin{aligned} \text{Neuron 3} & \begin{cases} w_{31} = w_{31} + \Delta w_{31} = 1 - 0.0043 = 0.9957 \\ w_{32} = w_{32} + \Delta w_{32} = -1 - 0.0150 = -1.0150 \\ w_{33} = w_{33} + \Delta w_{33} = 2 + 0.0177 = 2.0177 \end{cases} \\ \text{Neuron 4} & \begin{cases} w_{41} = w_{41} + \Delta w_{41} = 2 + 0.0788 = 2.0788 \\ w_{42} = w_{42} + \Delta w_{42} = -2 + 0.2730 = -1.7270 \\ w_{43} = w_{43} + \Delta w_{43} = 1 - 0.3218 = 0.6782 \end{cases} \end{aligned}$$



## EBP Error Back Propagation algorithm

### example

#### Algorithm steps:

- repeat the procedure for each pattern
- after applying all the patterns, start next iteration
- exit the algorithm once the TE is  $< \text{req.error}$

#### Problems:

- Oscillations
- Slow convergence or no convergence
- Flat spot, local minima traps

## Things to remember...

- **EBP introduced...**

- *Multi-layer, multi-output architecture*
- *mapping of  $n$ -dimensional to  $m$ -dimensional vector*
- *soft-activation function (historically)*

- **EBP represented...**

- *A go-to algorithm for a long time (still in use, new life with deep learning)*
- *Robust and reliable choice (though not a snake oil, e.g. slow/no convergence, oscillations, etc.)*