

Today

- **Threads vs Processes**
- Thread Features and types
- Hello world with threads
- Comparison and applications
- Concurrent Programming – Intro
- Sharing among threads



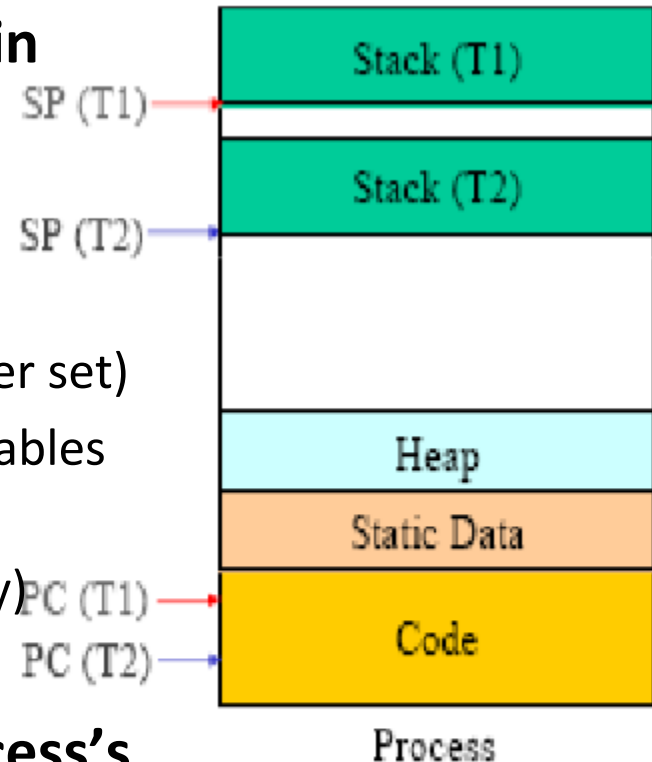
Motivation

- **Processes are expensive to create**
 - Address space for code and data pages
 - OS descriptors of resources allocated
 - State
 - Registers, etc ...
- **It takes quite a bit of time to switch between processes**
- **Communicating between processes is costly because most communication goes through the OS**
 - Files
 - Pipes
 - Message Queues
 - Shared memory



Threads

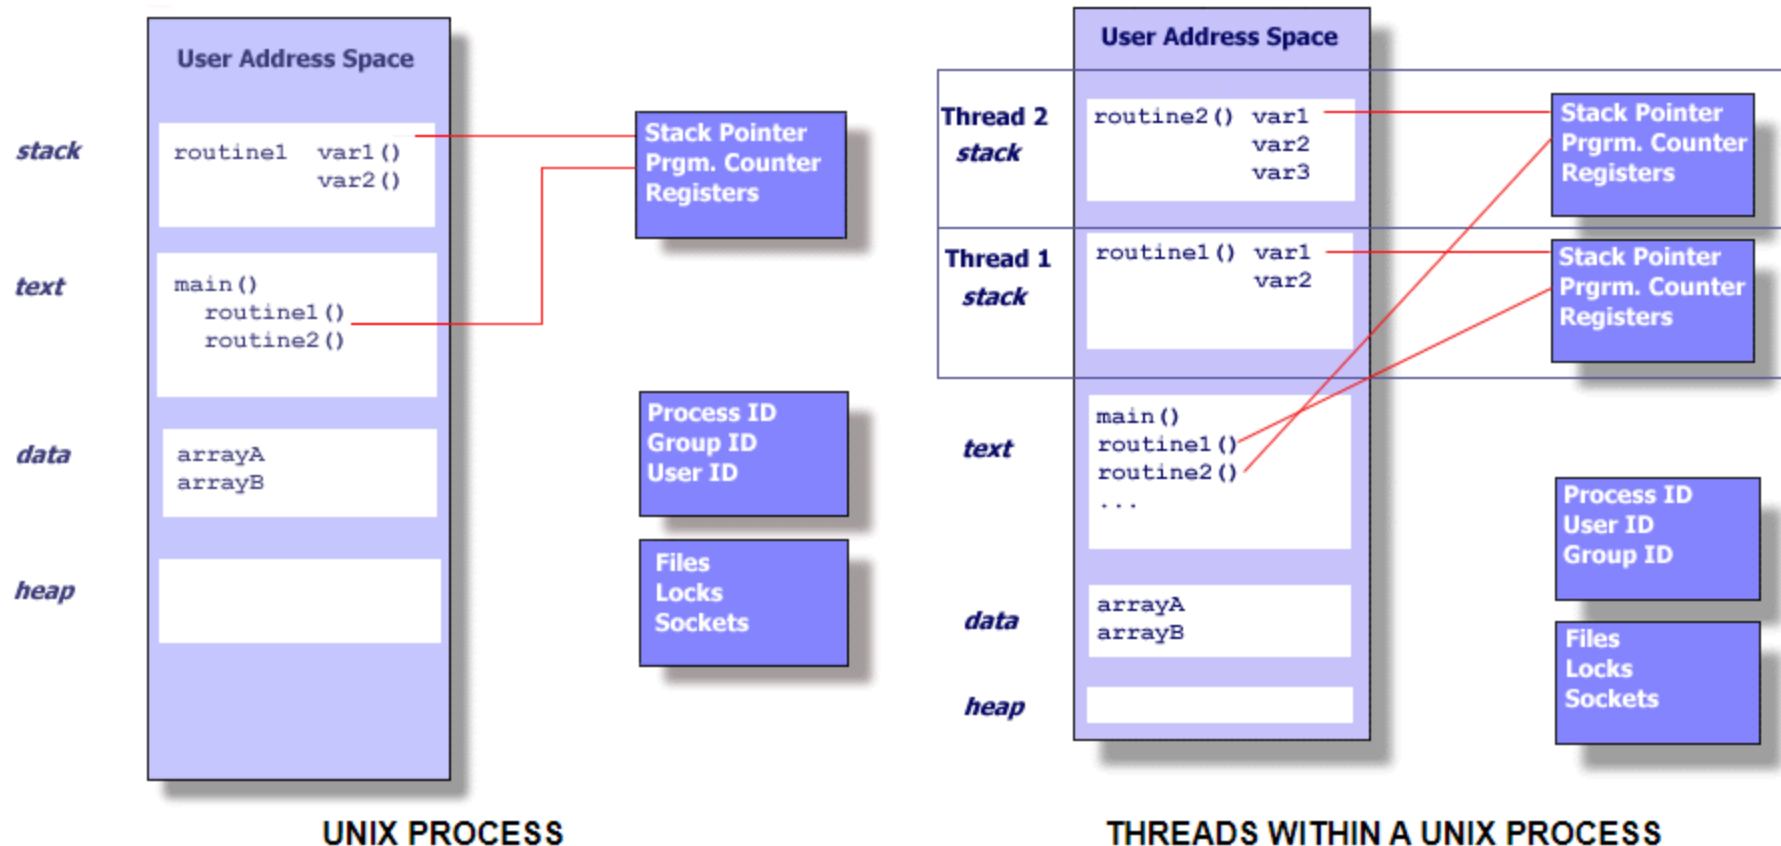
- A thread is the smallest schedulable unit in multithreading
- A thread in execution works with
 - Thread ID
 - Registers (program counter and working register set)
 - Stack (for procedure call parameters, local variables etc.)
 - Scheduling properties (such as policy or priority)
 - Set of pending and blocked signals
- A thread shares with other threads a process's (to which it belongs to)
 - Instruction and Data Sections
 - Permissions
 - Other resources, such as files



***Process with
2 threads***



Threads in the VM



<https://computing.llnl.gov/tutorials/pthreads/>



Threads vs. Processes

■ How threads and processes are similar



- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

■ How threads and processes are different

- Threads share all code and data (except local stacks)
 - Processes (typically) do not
- Threads are somewhat less expensive than processes
 - Process control (creating and reaping) twice as expensive as thread control
 - Linux numbers:
 - ~20K cycles to create and reap a process
 - ~10K cycles (or less) to create and reap a thread



Difference between Single vs Multithread Process



- A process by itself can be viewed a single thread and is traditionally known as a heavy weight process



Today

- Threads vs Processes
- **Thread Features and types, functions**
- Hello world with threads
- Comparison and applications
- Concurrent Programming – Intro
- Sharing among threads



Thread Features

1. **Exists within a process and uses the process resources**
2. **Has its own independent flow of control as long as its parent process exists and the OS supports it**
3. **Independently “schedulable”**
4. **May share the process resources with other threads**
5. **Dies if the parent process dies - or something similar**
6. **Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.**



Caveats

Because threads within the same process share resources:


- **Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.**
- **Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer.**
- **If one of the threads crash, whole process may crash**




Logical View of Threads

- **Threads associated with process form a pool of peers**
 - Unlike processes which form a tree hierarchy

Threads associated with process foo



Process hierarchy



Implementation

- Threads can be implemented in
 - *Kernel-space: Involves system calls, and library support*
 - *User-space: API functions implemented in user level*
- **POSIX PThreads** - may be provided as either a user or kernel library, as an extension to the POSIX standard.
- PThreads are available on Solaris, Linux, Mac OSX, Tru64 UNIX, and via public domain shareware for Windows.



Linux Threads – Kernel Level System Call

- Created through:
 - `clone (fn, stack, flags)`
- Where
 - `fn` specifies function to be executed by the new thread or process
 - `stack` points to the stack it will use
 - `flags` is a set of flags specifying various options (**CLONE_VM** or **CLONE_THREAD** for threads. If it is missing a regular process will be created)
 - `fork()` also calls `clone` internally
 - Threads are also called lightweight processes in kernel space .



Linux Threads – User Level with Kernel Support

- POSIX threads, or ***PThreads***, started as pure user-level threads managed by the POSIX thread library
- IEEE standard, supported by most vendors
 - Emerged as standard thread API
- Gained later ***some kernel support***
- Ported to various Unix and Windows systems
- Function names start with **pthread_**
- Calls tend to have a complex syntax



Posix Threads (Pthreads) Interface

- ***Pthreads***: Standard interface for ~60 functions that manipulate threads from C programs
 - Creating and reaping threads
 - `pthread_create()`
 - `pthread_join()`
 - Determining your thread ID
 - `pthread_self()`
 - Terminating threads
 - `pthread_cancel()`
 - `pthread_exit()`
 - `exit()` [terminates all threads]
 - Synchronizing access to shared variables (we will use semaphores instead)
 - `pthread_mutex_init`
 - `pthread_mutex_[un]lock`




Thread lifecycle

Main

- ✓ `pthread_create()`
(create thread)
- ✓ wait for thread() to finish via
`pthread_join()`

Thread

- ✓ begins at function pointer
- ✓ runs until
`pthread_exit()`



pthread_create()

- The pthread_create function starts a new thread in the calling process.

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine) (void *),  
                  void *arg);
```

- Where,
 - ▶ **thread** is a pthread library structure holding thread info
 - ▶ **attr** is a set of attributes to apply to the thread (Usually NULL)
 - ▶ **start_routine** is the thread function pointer
 - ▶ **arg** is an opaque data pointer to pass to thread
- Returns 0 on success



pthread_join()

- The pthread_join function waits for the thread specified by thread to terminate.

```
int pthread_join(pthread_t thread, void **retval);
```

- Where,
 - ▶ **thread** is a pthread library structure holding thread info
 - ▶ **retval** is a double pointer return value. Usually NULL if not will return the exit status
- Will return 0 if successful
- Similar to waitpid for fork. There is no wait alternative for Pthreads



pthread_exit()

- The pthread_exit function terminates the calling thread and returns a value

```
void pthread_exit(void *retval);
```

- Where,
 - ▶ **retval** is a pointer to a return value
- Note:** better be dynamically allocated because the thread stack will go away when the thread exits



Today

- Threads vs Processes
- Thread Features and types
- **Hello world with threads**
- Comparison and applications
- Concurrent Programming – Intro
- Sharing among threads



The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

Thread ID

Thread attributes
(usually NULL)

Thread routine

Thread arguments
(void *p)

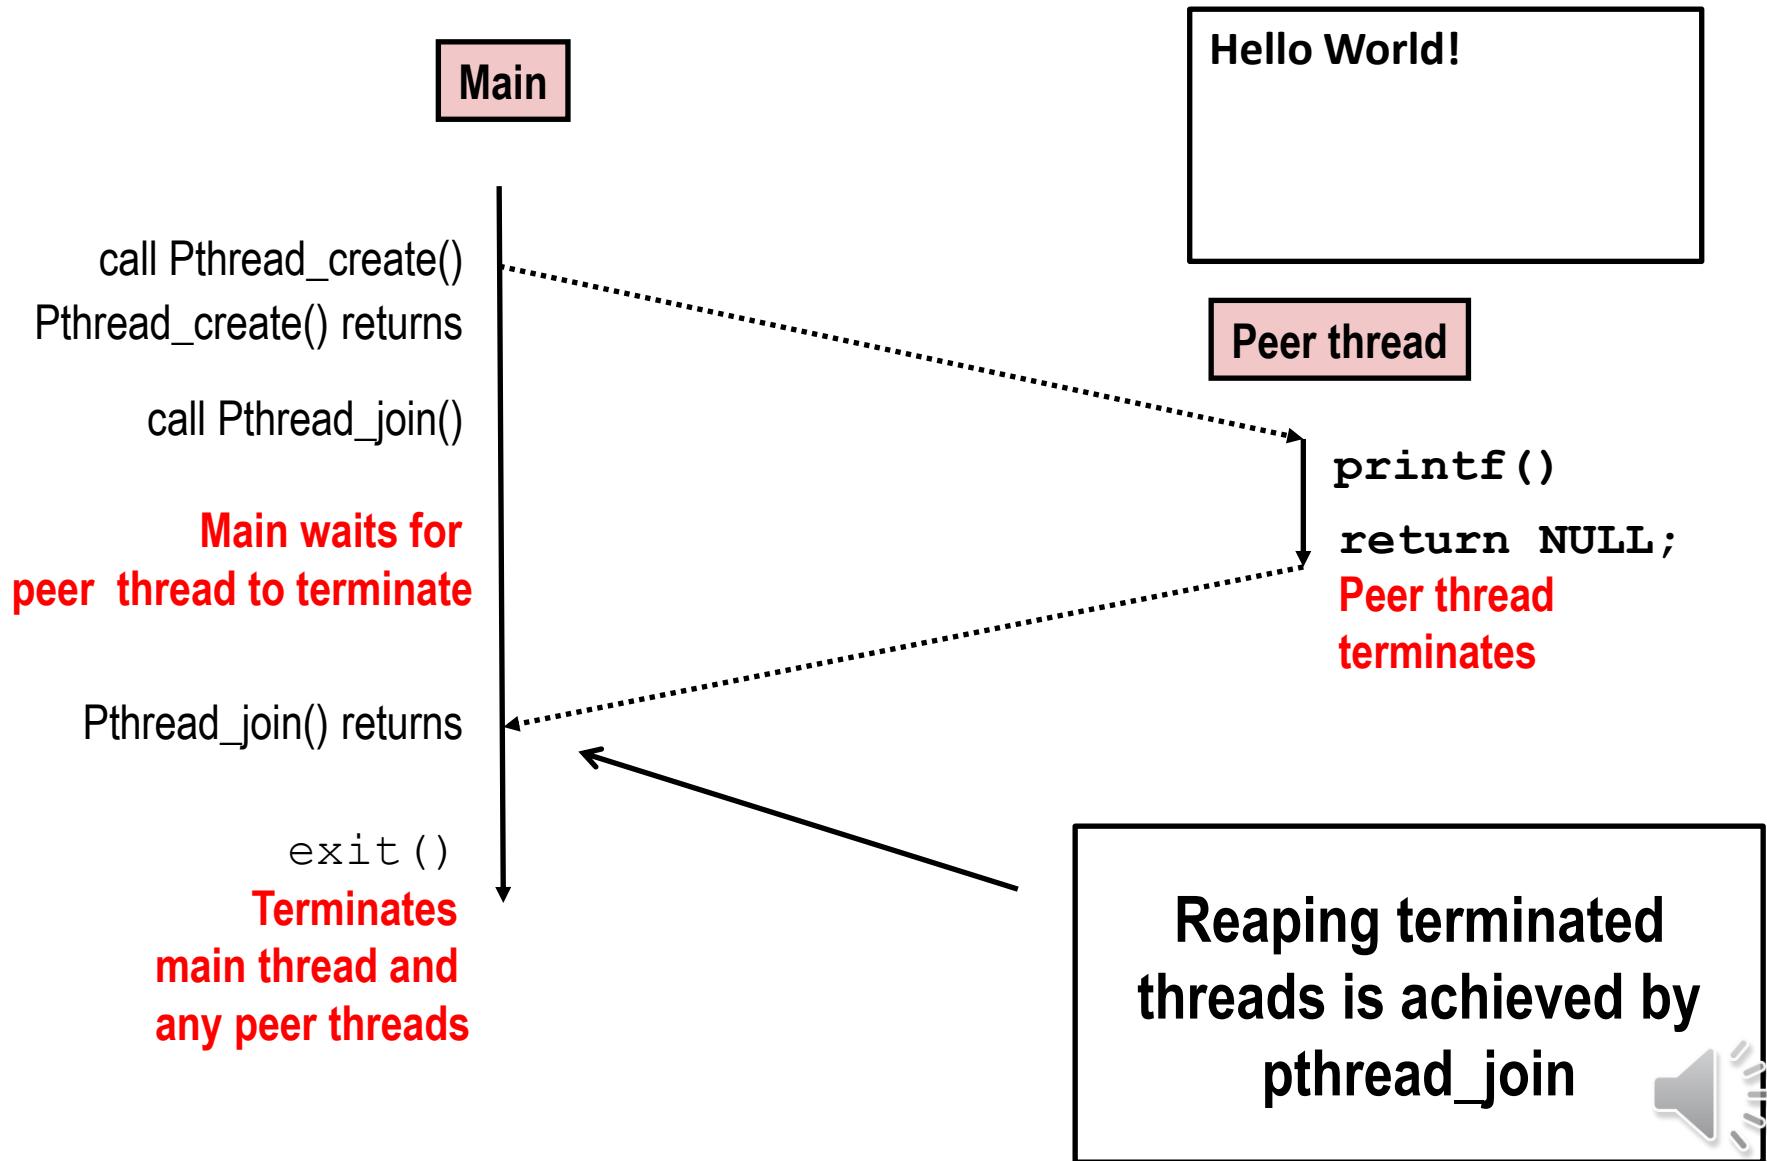
Return value
(void **p)

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c



Execution of Threaded “hello, world”



Putting it all together ...

```
typedef struct {
    int num;
    const char *str;
} MY_STRUCT;

void * thread_function( void * arg ) {
    MY_STRUCT *val = (MY_STRUCT *)arg; // Cast to expected type
    printf( "Thread %lx has vaules %x,%s]\n", pthread_self(), val->num, val->str );
    pthread_exit( &val->num );
}

int main( void ) {
    MY_STRUCT v1 = { 0x12345, "Val 1" };
    MY_STRUCT v2 = { 0x54312, "Val 2" };
    pthread_t t1, t2;
    printf( "Starting threads\n" );
    pthread_create( &t1, NULL, thread_function, (void *)&v1 );
    pthread_create( &t2, NULL, thread_function, (void *)&v2 );
    pthread_join( t1, NULL );
    pthread_join( t2, NULL );
    printf( "All threads returned\n" );
    return( 0 );
}
```




Putting it all together ...

```
typedef struct {
    int num;
    const char *str;
} MY_STRUCT;

void * thread_function( void * arg ) {
    MY_STRUCT *val = (MY_STRUCT *)arg; // Cast to expected type
    printf( "Thread %lx has vaules %x,%s]\n", pthread_self(), val->num, val->str );
    pthread_exit( &val->num );
}

int main( void ) {
    MY_STRUCT v1 = { 0x12345, "Val 1" };
    MY_STRUCT v2 = { 0x54312, "Val 2" };
    pthread_t t1, t2;
    printf( "Starting threads\n" );
    pthread_create( &t1, NULL, thread_function, (void *)&v1 );
    pthread_create( &t2, NULL, thread_function, (void *)&v2 );
    pthread_join( t1, NULL );
    pthread_join( t2, NULL );
    printf( "All threads returned\n" );
    return( 0 );
}
```

```
$ ./concurrency
Starting threads
Thread 7f51c3e05700 has vaules 54312,Val 2]
Thread 7f51c4606700 has vaules 12345,Val 1]
All threads returned
$
```



Today

- Threads vs Processes
- Thread Features and types
- Hello world with threads
- **Comparison and applications**
- Concurrent Programming – Intro
- Sharing among threads



When to use processes

- **When fault tolerance needed: If child process crashes parent process is unaffected. Heavily used in web-servers & databases to improve resilience**
 - If a single thread crashes, it may crash all other threads due to shared memory
- **Security: Parent and child can have different security settings, can run as different users.**
 - Threads are designed to share.



When to use threads

- **When sharing and/or collaboration needed**
- **For performance critical applications**
 - To gain speed through in multi-processor, multi-core systems. In multiprocessor system threads are more effective.
- **If security is not critical**
 - Buffer overflow will enable accessing to the stack of other threads



Today

- Threads vs Processes
- Thread Features and types
- Hello world with threads
- Comparison and applications
- **Concurrent Programming – Intro**
- Sharing among threads



Concurrent Programming is Hard!

- The human mind tends to be sequential
- The notion of time is often misleading
- Thinking about all possible sequences of events in a computer system is at least error prone and frequently impossible



Concurrent Programming is Hard!

- **Classical problems of concurrent programs:**

- **Races:** outcome depends on arbitrary scheduling decisions elsewhere in the system. There is no natural way of knowing if another thread is already working with same data.

- **Requires mutual exclusion**

- **Hardware Level Instructions**

- test-and-set, compare_and_swap

- **OS support**

- Mutex Locks, Semaphores

- **Higher level software support**


- Monitors



Race condition example on a Bounded Buffer

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
while (true) {  
    while (counter == 0); /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```



Race Condition

- **counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = counter
S1: producer execute register1 = register1 + 1
S2: consumer execute register2 = counter
S3: consumer execute register2 = register2 - 1
S4: producer execute counter = register1
S5: consumer execute counter = register2
```

Assume initial value of
counter = 5

- Adding one to the buffer
- Removing one from buffer

register1 = 5

register1 = 6

register2 = 5

register2 = 4

counter = 6

counter = 4

Any other possible results?



Race Condition

- We would arrive this incorrect state because we allowed both processes to manipulate the variable counter concurrently.
- A situation like this is called race condition
- We need to ensure that only one at a time can manipulate variable counter
 - To make such a guarantee, we need to synchronize processes in some way
- Synchronization Brings other problems on the table



Concurrent Programming is Hard!

- **Synchronization brings other problems to the table:**
 - **Deadlock:** improper resource allocation prevents forward progress
 - Example: traffic gridlock
 - **Starvation / Fairness:** external events and/or system scheduling decisions can prevent sub-task progress
 - Example: people always jump in front of you in line

