



Lecture Outline

- **C's Memory Model**
- **Pointers (an introduction)**
- **Arrays**



OS and Processes

- **The OS lets you run multiple applications at once**
 - An application runs within an OS “process”
 - The OS time slices each CPU between runnable processes
 - This happens *very quickly*: ~100 times per second



Processes and Virtual Memory (VM)

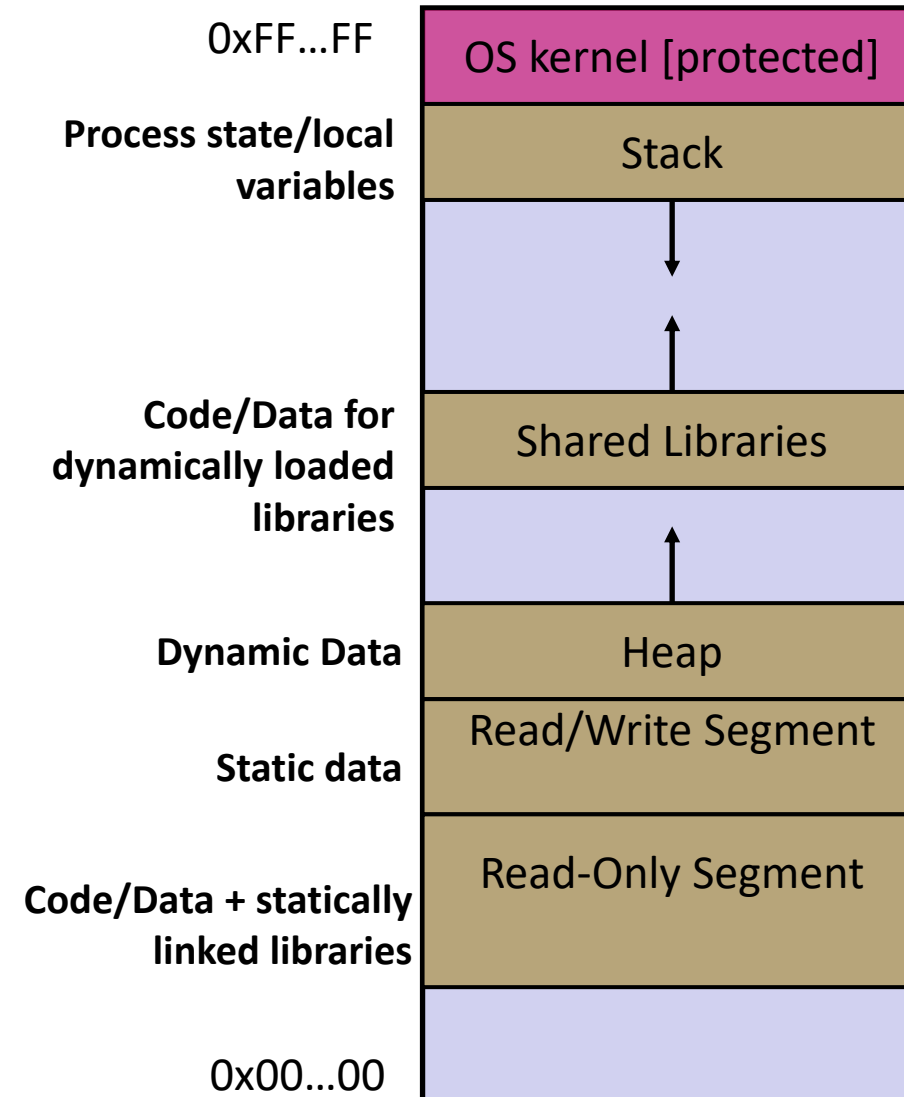
- The OS gives each process the illusion of its own private memory
 - Called the process' **address space**
 - Contains the process' virtual memory, visible only to it (via translation)
 - 2^{64} bytes on a 64-bit machine



Loading Process into VM

■ When the OS loads a program it:

- 1) Creates an address space
- 2) Inspects the executable file to see what's in it
- 3) (Lazily) copies regions of the file into the right place in the address space
- 4) Does any final linking, relocation, or other needed preparation



Process Memory

■ *Local* variables on the Stack


- Allocated and freed via calling conventions (`push`, `pop`, `mov`)

■ *Global* and *static* variables in Data

- Allocated/freed when the process starts/exits

■ *Dynamically-allocated* data on the Heap

- `malloc()` to request; `free()` to free, otherwise **memory leak**



Note on Static Variables and Functions in C

■ Static variables:

- Preserve their previous value in their previous scope and are not initialized again in the new scope.
- Static variables and global variables are initialized to 0 if not explicitly initialized

```
#include<stdio.h>
int fun()
{
    static int count = 0;
    count++;
    return count;
}

int main()
{
    printf("%d ", fun());
    printf("%d ", fun());
    return 0;
}
```

Output will be:

1 2

<https://ide.geeksforgeeks.org/K2eSmAECU1>

Link to run on online
IDE

Note on Static Variables and Functions in C

- Functions are global by default. The “static” keyword before a function name makes it static
- Access to static functions is restricted to the file where they are declared

```
/* Inside file1.c */
static void fun1(void)
{
    puts("fun1 called");
}
```

```
/* Inside file2.c */
int main(void)
{
    fun1();
    getchar();
    return 0;
}
```

- if we compile the above code with command “*gcc file2.c file1.c*”, we get the error “*undefined reference to `fun1`*”. This is because *fun1()* is declared *static* in *file1.c* and cannot be used in *file2.c*.

How about stack and stack frame?

■ Why stacks?

- To be able to restore previous state of the caller function
- To store information that cannot fit into the limited number of registers of the CPU

■ Uses of stack

- Pass function arguments, Store return information
- Save registers for later restoration,
- Local variables.

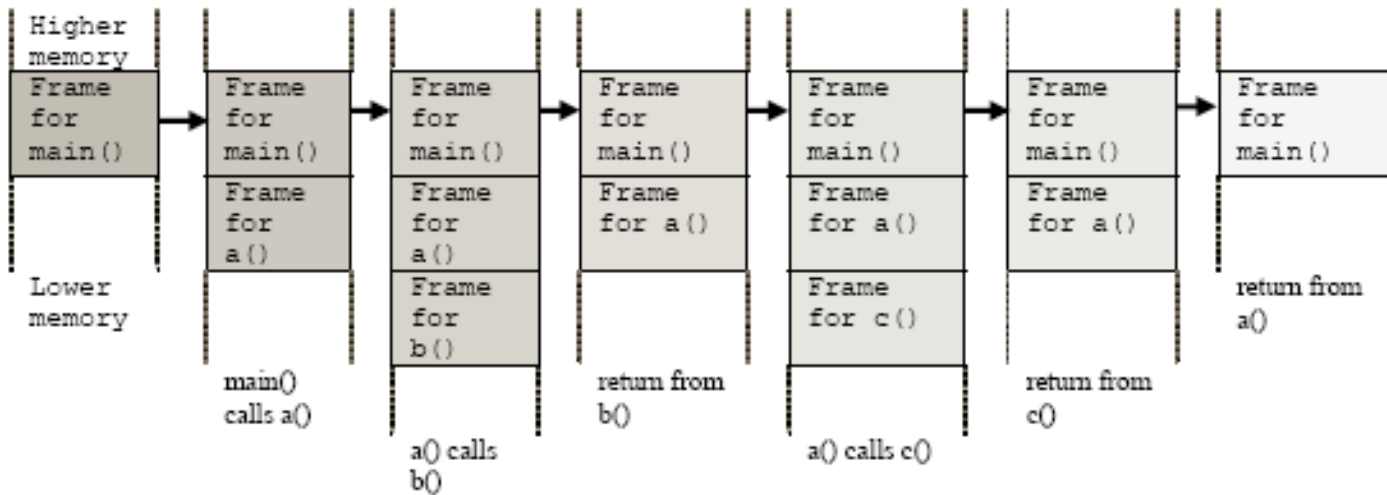
■ The portion of the stack allocated for a single function call is called a stack frame aka activation record.

- allocated when a function is called,
- de-allocated when it returns.

Stack frame?

- For each function call, a new stack frame is created on the stack.

- allocated when a function is called,
- de-allocated when it returns.



```
#include <stdio.h>
```

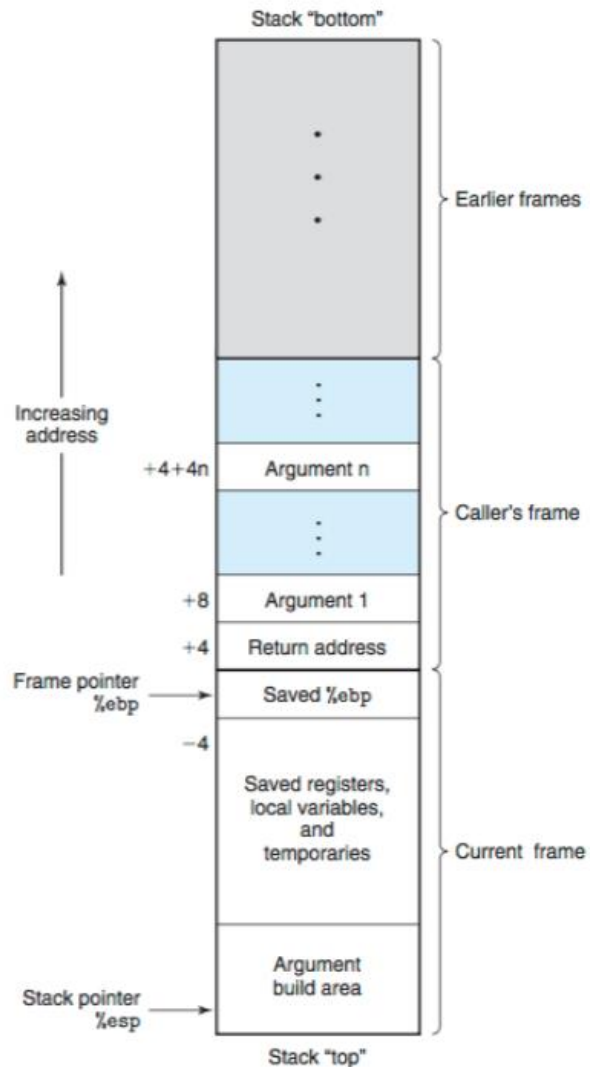
```
int b()  
{ return 0; }
```

```
int c()  
{ return 0; }
```

```
int a()  
{  
    b();  
    c();  
    return 0;  
}
```

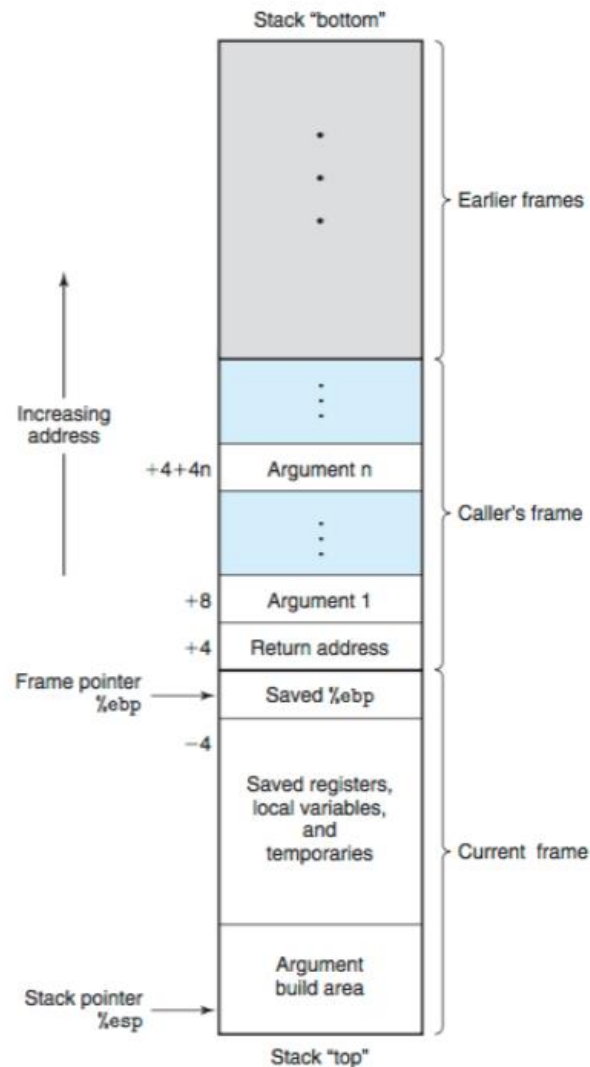
```
int main()  
{  
    a();  
    return 0;  
}
```

What is in the stack frame



- Before executing a function, all of the **arguments** for the function are pushed onto the stack in the reverse order
- the address of the next instruction within the caller function, which is the **return address (previous %eip)**, is pushed onto the stack.
- instruction pointer %eip is updated to point to the start of the function.

What is in the stack frame



- The first instruction is to save the current base pointer register **%ebp**. (Which belongs to the caller function)
- **%ebp** serves as the reference point for the stack frame
- Once the current function is done, we need to resume the execution of the caller function. This means that we need to restore the caller's base pointer register **%ebp**. Thus, we need to save the caller's base pointer register (**Saved %ebp**).
- Stack pointer, **%esp** is a pointer to the top of the stack

Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c

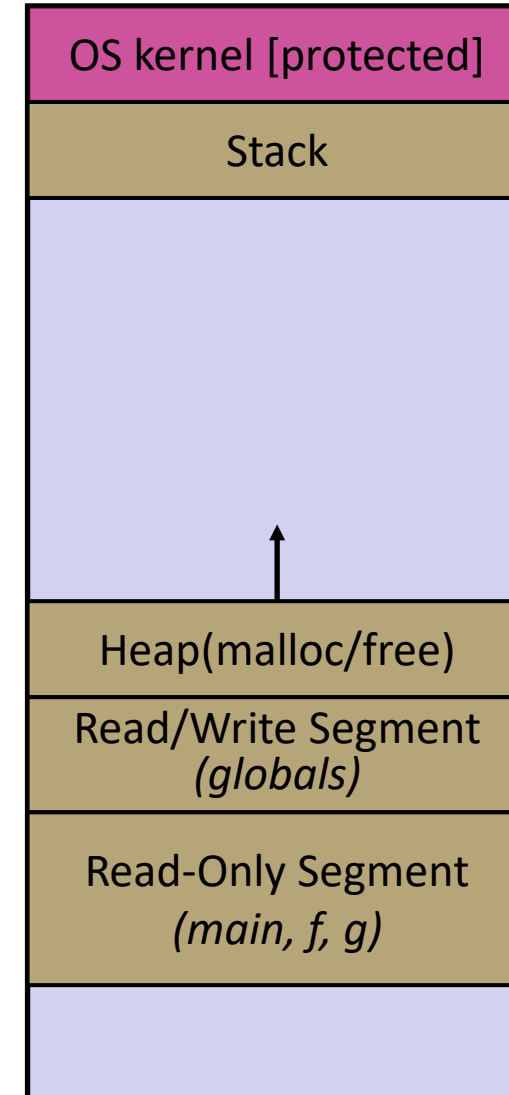
```
#include <stdint.h>

int f(int, int);
int g(int);

→ int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c




```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

→ int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed (like in gdb).

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c

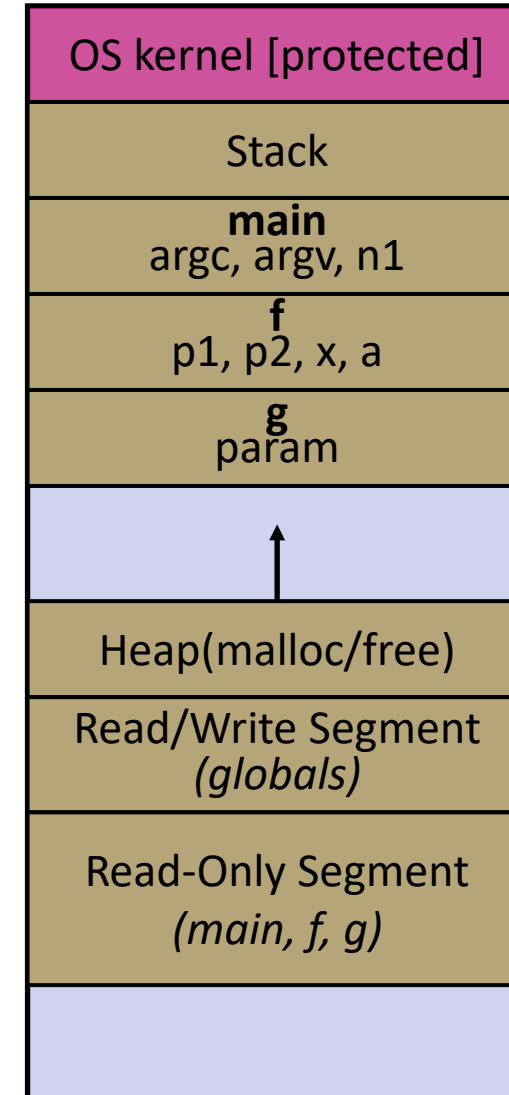
```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c

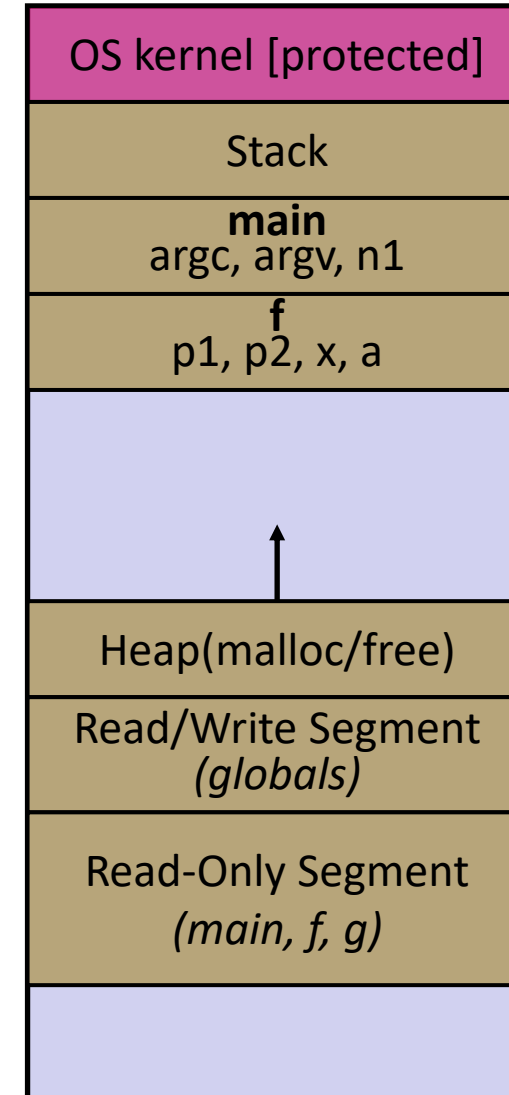
```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

→ int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c

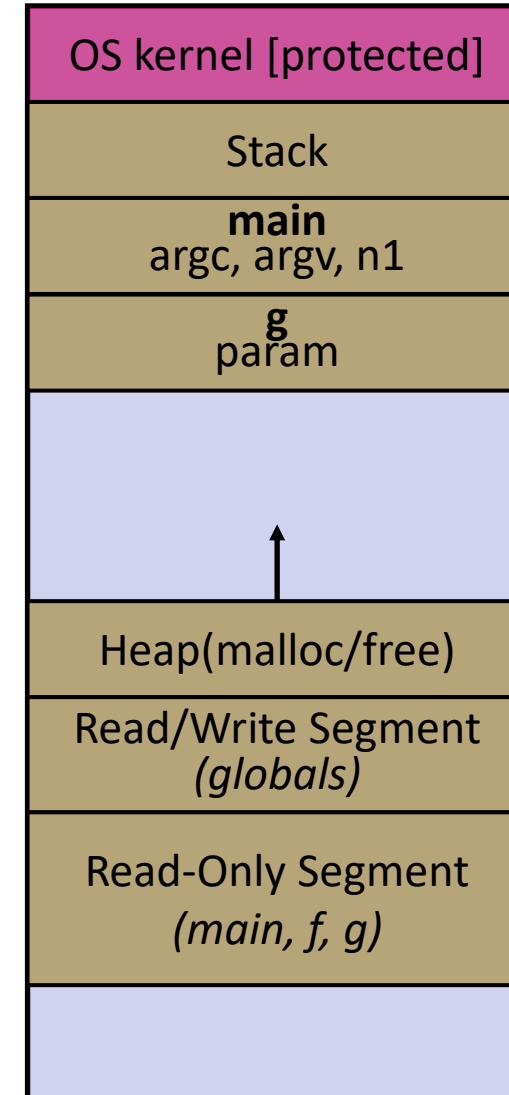
```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action

Note: arrow points to *next* instruction to be executed

stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Lecture Outline

- C's Memory Model
- Pointers (an introduction)
- Arrays

Pointers

■ Variables that store addresses

- Pointers points to somewhere in the process' virtual address space. (contain address withing VM)

■ Declaration: `type* ptr; or type *ptr;`

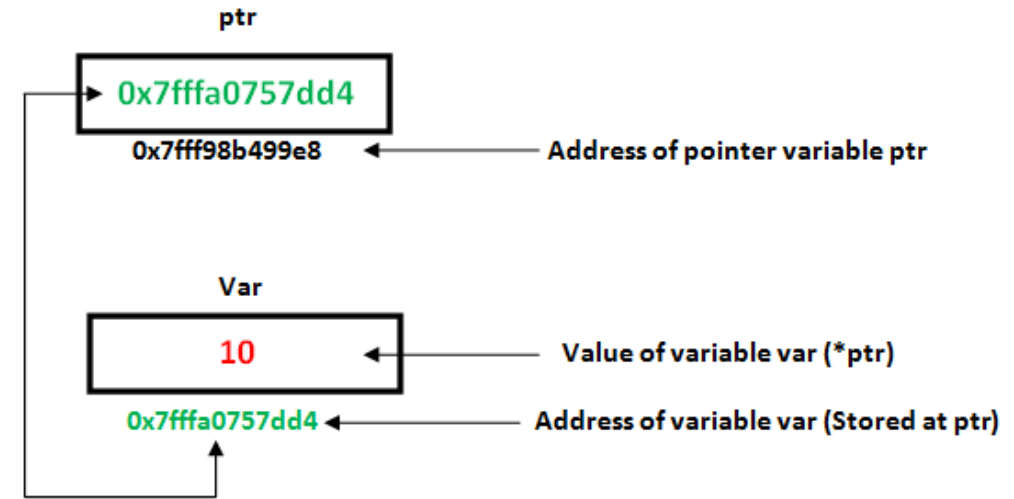
- Recommended: do not declare multiple pointers on same line:
`int *p1, p2` or `int* p1, p2 ;`

- Are not the same as `int *p1, *p2;`

- Instead, use: `int *p1;`
`int *p2;` or `int *p1;`
`int *p2;`

- `&var` produces the virtual address of `var`

- `ptr = &var,` stores the address of `var` into `ptr`



- *Dereference* a pointer using the unary `*` operator
 - Access the memory referred to by a pointer
 - `*ptr` is equivalent of `var`

Pointer Example

pointy.c

```
#include <stdio.h>
#include <stdint.h>

int main(int argc, char** argv) {
    int x = 351;
    int* p;      // p will be a pointer to a int

    p = &x;      // p now contains the addr of x
    printf("&x is %p\n", &x);
    printf(" p is %p\n", p);
    printf(" x is %d\n", x);

    *p = 333;    // change the value of x
    printf(" x is %d\n", x);

    return 0;
}
```

```
&x is 0x7ffda992e044
p is 0x7ffda992e044
x is 351
x is 333
```

To practice online:

<https://onlinegdb.com/HkwuOW9bl>

Something Curious

- What happens if we run `pointy.c` several times?

```
$ gcc -Wall -std=c11 pointy.c -o pointy
```

Run 1:

```
$ ./pointy
&x is 0x7ffff9e28524
p is 0x7ffff9e28524
x is 351
x is 333
```

Run 2:

```
$ ./pointy
&x is 0x7fffe847be34
p is 0x7fffe847be34
x is 351
x is 333
```

Run 3:

```
$ ./pointy
&x is 0x7fffe7b14644
p is 0x7fffe7b14644
x is 351
x is 333
```


Run 4:

```
$ ./pointy
&x is 0x7fffff0dfe54
p is 0x7fffff0dfe54
x is 351
x is 333
```

Address Space Layout Randomization

- Linux uses *address space layout randomization (ASLR)* for added security

- Randomizes:
 - Base of stack
 - Shared library (`mmap`) location
- Makes Stack-based buffer overflow attacks tougher
- Makes debugging tougher
- Can be disabled (`gdb` does this by default); Google if curious



Lecture Outline

- C's Memory Model
- Pointers (An Introduction)
- Arrays

Arrays

■ Definition:

```
type name [size]
```

- Allocates `size*sizeof(type)` bytes of *contiguous* memory
- Normal usage is a compile-time constant for `size` (e.g. `int scores[175];`)
- Initially, array values are “garbage”

■ Size of an array

- Not stored anywhere – array does not know its own size!
 - `sizeof(array)` only works in variable scope of array definition
 - If you are passing arrays as a parameter to a function, it decays to a pointer.
- Recent versions of C (but *not* C++) allow for variable-length arrays
 - Uncommon and can be considered bad practice (Linux Kernel doesn't use it, Look what Linus Torvalds say): https://en.wikipedia.org/wiki/Variable-length_array

```
int n = 175;  
int scores[n]; // OK in C99
```

**sizeof() operator
in C calculates the
size of its operand**

Using sizeof () with arrays

```
1  #include <stdio.h>
2
3  void foo();
4
5  int main(int argc, char** argv) {
6      int numbers[] = {9, 8, 1, 9, 5};
7      printf("Size of a single integer %lu\n", sizeof(numbers[0]));
8      printf("Size of integer type %lu\n", sizeof(int));
9      printf("Size of address of an integer %lu\n", sizeof(&numbers[0]));
10     printf("Size of numbers within main (caller) function %lu\n", sizeof(numbers));
11     foo(numbers);
12     return 0;
13 }
14
15 void foo(int* numbers) {
16     printf("Size of numbers within foo (callee) function: %lu", sizeof(numbers));
17 }
```

Runnable link:

<https://onlinegdb.com/BJ-pqEWs8>

Size of a single integer 4

Size of integer type 4

Size of address of an integer 8

Size of numbers within main (caller) function 20

Size of numbers within foo (callee) function: 8

Using Arrays

■ **Initialization:** `type name[size] = {val0, ..., valN};`

- `{ }` initialization can *only* be used at time of definition
- If no `size` supplied, infers from length of array initializer

■ **Array name used as identifier for “collection of data”**

- `name[index]` specifies an element of the array and can be used as an assignment target or as a value in an expression
- Array name (by itself) produces the address of the start of the array

```
int primes[6] = {2, 3, 5, 6, 11, 13};  
primes[3] = 7;  
primes[100] = 0; // memory smash!
```

Special cases for initialization

- ❑ If fewer values are given than the size of the array, empty indices are filled with zero

`int primes[3] = {2};` is equivalent to `int primes[3] = {2, 0, 0};`

- ❑ If more values are given than the size of the array, values that don't fit are ignored and the compiler gives a warning: excess elements in array initializer

`int primes[3] = {2, 3, 4, 5};` //will cause compile warning

- ❑ For one dimensional arrays, providing the size is optional

`int primes[] = {2, 3, 4};`

same as

`int primes[3] = {2, 3, 4};`

Exercise 1

- What is the output of the following program?

```
int main()
{
    int i;
    int arr[5] = {1};
    for (i = 0; i < 5; i++)
        printf("%d ", arr[i]);
    return 0;
}
```

- A) 1 followed by four garbage values
- B) 1 0 0 0 0
- C) 1 1 1 1 1
- D) 0 0 0 0 0

Exercise 2

- What is the output of the following program?

```
int main()
{
    int i;
    int arr[5] = {0};
    for (i = 0; i <= 5; i++)
        printf("%d ", arr[i]);
    return 0;
}
```

- A) Compiler Error: Array index out of bound.
- B) Always Prints 0 five times followed by garbage value
- C) The program always crashes.
- D) The program may print 0 five times followed by garbage value, or may crash if address (arr+5) is invalid.

Exercise 3

- ## ■ What is the output of the following program?

```
#include <stdio.h>

int main()
{
    int arr[50] = {0,1,2,[47]=47,48,49};

    for(int i=0 ; i<50 ;i++)
        printf ("%d ", arr[i]);
}
```

- In C, initialization of array can be done for selected elements as well. By default, the initializer start from 0th element. Specific elements in array can be specified by []

Program will initialize arr[0], arr[1], arr[2], arr[47], arr[48] and arr[49] to 0,1,2,47,48 and 49 respectively. The remaining elements of the array would be initialized to 0.

[illegible]

Multi-dimensional Arrays

■ Generic 2D format:

```
type name[rows][cols] = {{values}, ..., {values}};
```

- Still allocates a single, contiguous chunk of memory
- C is *row-major*

```
// a 2-row, 3-column array of doubles
double grid[2][3];

// a 3-row, 5-column array of ints
int matrix[3][5] = {
    {0, 1, 2, 3, 4},
    {0, 2, 4, 6, 8},
    {1, 3, 5, 7, 9}
};
```

Every dimension beyond the first must be specified!

Exercise 4

- What is the output of the following program?

```
#include <stdio.h>

int main()
{
    int a[][] = {{1,2},{3,4}};
    int i, j;
    for (i = 0; i < 2; i++)
        for (j = 0; j < 2; j++)
            printf("%d ", a[i][j]);
    return 0;
}
```

- A) 1 2 3 4
- B) Compiler Error in line " int a[][] = {{1,2},{3,4}};"
- C) 4 garbage values
- D) 4 3 2 1

Exercise 5

■ Valid or invalid?

/ Valid declaration*/*

int abc[2][2] = {1, 2, 3 ,4 }

/ Valid declaration*/*

int abc[][2] = {1, 2, 3 ,4 }

/ Invalid declaration – you must specify second dimension*/*

int abc[][] = {1, 2, 3 ,4 }

/ Invalid because of the same reason mentioned above*/*

int abc[2][] = {1, 2, 3 ,4 }

Parameters: reference vs. value

- There are two fundamental parameter-passing schemes in programming languages
- **Call-by-value**
 - Parameter is a local variable initialized with a copy of the calling argument when the function is called; manipulating the parameter **only changes the copy**, *not* the calling argument
 - C, Java, C++ (most things)
- **Call-by-reference**
 - Parameter is an alias for the supplied argument; manipulating the parameter **manipulates the calling argument**
 - C++ references (Actually is a C++ thing but simulated in C)

Arrays as Parameters

■ It's tricky to use arrays as parameters

- What happens when you use an array name as an argument?
- Arrays do not know their own size

```
int sumAll(int a[]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    return 0;
}

int sumAll(int a[]) {
    int i, sum = 0;
    for (i = 0; i < ...???)
}
```

Solution 1: Declare Array Size

```
int sumAll(int a[5]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    printf("sum is: %d\n", sum);
    return 0;
}

int sumAll(int a[5]) {
    int i, sum = 0;
    for (i = 0; i < 5; i++) {
        sum += a[i];
    }
    return sum;
}
```

- Problem: loss of generality/flexibility


<https://onlinegdb.com/ByL93WqW8>

Solution 2: Pass Size as Parameter

```
int sumAll(int a[], int size); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers, 5);
    printf("sum is: %d\n", sum);
    return 0;
}

int sumAll(int a[], int size) {
    int i, sum = 0;
    for (i = 0; i < size; i++) {
        sum += a[i];
    }
    return sum;
}
```

A red arrow originates from the number '5' in the function call `sumAll(numbers, 5)` within the `main` function. It points downwards and to the right, ending at the `size` parameter in the function definition `sumAll(int a[], int size)`. This illustrates how the size of the array is passed as an argument to the function.

arraysum.c

- This is the standard idiom in C programs

<https://onlinegdb.com/HkwMa-q-U>

Returning an Array

- **Local variables, including arrays, are allocated on the Stack**
 - They “disappear” when a function returns!
 - Can’t safely return local arrays from functions
 - Can’t return an array as a return value

```
int* copyArray(int src[], int size) {  
    int i, dst[size];    // OK in C99  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
  
    return dst;    // no compiler error, but wrong!  
}
```

buggy_copyarray.c

Solution: Output Parameter

- Create the “returned” array in the caller
 - Pass it as an **output parameter** to `copyarray()`
 - A pointer parameter that allows the called function to store values that the caller can use
 - Works because arrays are “passed” as pointers
 - “Feels” like call-by-reference, *but technically it's not*

```
void copyArray(int src[], int dst[], int size) {  
    int i;  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
}
```

`copyarray.c`

<https://onlinegdb.com/qjFnC-4lAt>

Arrays: Call-By-Value or Call-By-Reference?

- **Technical answer:** a $\mathbb{T} []$ array parameter is “promoted” to a pointer of type \mathbb{T}^* , and the *pointer* is passed by value
 - So, it acts like a call-by-reference array (if callee changes the array parameter elements it changes the caller’s array)
 - But it’s really a call-by-value pointer (the callee can change the pointer parameter to point to something else(!))

```
void copyArray(int src[], int dst[], int size) {  
    int i;  
    dst = src;  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];    // copies source array to itself!  
    }  
}
```


Next

- Debugging
- Pointers