

Fork, Wait, Exec

- System call error handling
- More on fork System Call and Process Graphs
- wait and waitpid system calls
- exec family system calls
- signals

System Call Error Handling

- On error, Linux system-level functions typically return -1 and set global variable `errno` to indicate cause.
- Hard and fast rule:
 - You must check the return status of every system-level function
 - Only exception is the handful of functions that return `void`
- Example:

```
if ((pid = fork()) < 0) {  
    fprintf(stderr, "fork error: %s\n", strerror(errno));  
    exit(0);  
}
```

Error-reporting functions

- Can simplify somewhat using an *error-reporting function*:

```
void unix_error(char *msg) /* Unix-style error */  
{  
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));  
    exit(0);  
}
```

```
if ((pid = fork()) < 0)  
    unix_error("fork error");
```

Error-handling Wrappers

- Textbook simplifies the code further by using Stevens-style error-handling wrappers:

```
pid_t Fork(void)
{
    pid_t pid;

    if ((pid = fork()) < 0)
        unix_error("Fork error");
    return pid;
}
```

```
pid = Fork();
```

Note on sample programs

- Most code examples I use in this chapter are from text
- You can find them at :
<http://csapp.cs.cmu.edu/3e/code.html>
 - Under efc
- There are some custom libraries book is using to get them work you need to include csapp.h Link with csapp.c and use `-pthread` option.

csapp.cs.cmu.edu/3e/code.html

```
$ gcc -pthread fork.c csapp.c -o fork
```

```
$ ./fork
```

```
parent: x=0  
child : x=2
```



data/

- [show-bytes.c](#) [chap 2 (1 ref)]

ecf/

- [counterprob.c](#) [chap 8 (1 ref)]
- [fork.c](#) [chap 8 (1 ref)]
- [forkprob0.c](#) [chap 8 (1 ref)]
- [forkprob1.c](#) [chap 8 (1 ref)]
- [forkprob2.c](#) [chap 8 (1 ref)]

Today

- System call error handling
- More on fork System Call and Process Graphs
- wait and waitpid system calls
- exec family system calls
- signals

fork facts

■ Concurrent execution

- Can't predict execution order of parent and child

■ Duplicate but separate address space

- `x` has a value of 1 when `fork` returns in parent and child
- Subsequent changes to `x` are independent

■ Shared open files

- `stdout` is the same in both parent and child

Modeling fork with Process Graphs


- **A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program:**
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens before b
 - Edges can be labeled with current value of variables
 - `printf` vertices can be labeled with output
 - Each graph begins with a vertex with no inedges
- **Any *topological sort* of the graph corresponds to a feasible total ordering.**
 - Total ordering of vertices where all edges point from left to right

Process Graph Example

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

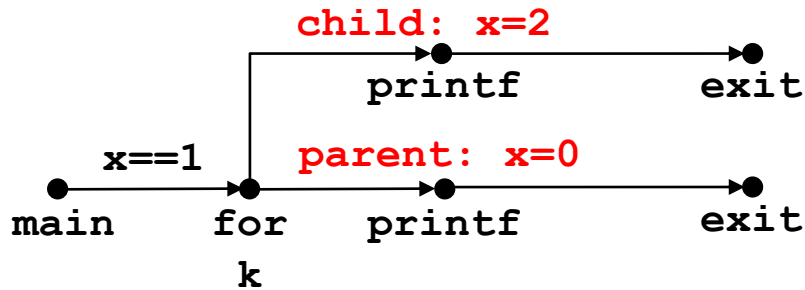
    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```




fork.c

Interpreting Process Graphs


■ Original graph:




■ Re-labelled graph:



Feasible total ordering:



Infeasible total ordering:




Practice

- How many lines of output would the program generate

(a) Calls fork once

```
1  #include "csapp.h"
2
3  int main()
4  {
5      Fork();
6      printf("hello\n");
7      exit(0);
8  }
```




Practice

- How many lines of output would the program generate

(c) Calls fork twice

```
1  #include "csapp.h"
2
3  int main()
4  {
5      Fork();
6      Fork();
7      printf("hello\n");
8      exit(0);
9  }
```




Practice

- How many lines of output would the program generate

(e) Calls fork three times

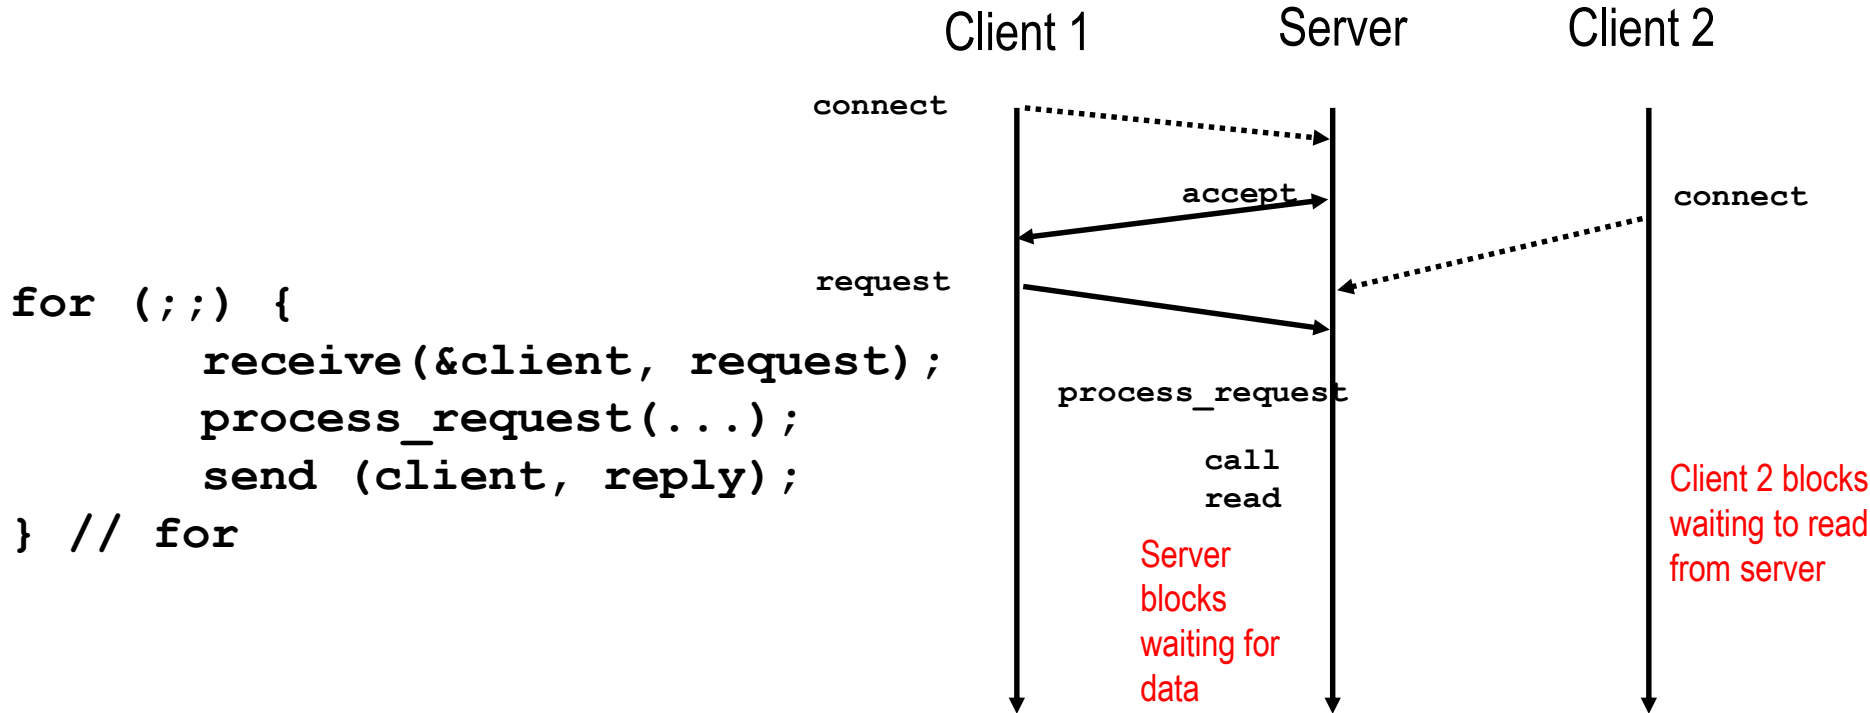
```
1  #include "csapp.h"
2
3  int main()
4  {
5      Fork();
6      Fork();
7      Fork();
8      printf("hello\n");
9      exit(0);
10 }
```



An application – Servers

- Most client requests involve disk accesses
 - File servers
 - Authentications servers
- When this happens, the server remains in the **BLOCKED** state if the server is iterative
- Cannot handle other customers' requests
- Analogy
 - *A waitperson that would only be able to wait on one table at a time would be idle most of the time.*

Fundamental Flaw of Iterative Servers



■ Solution: use *concurrent servers* instead

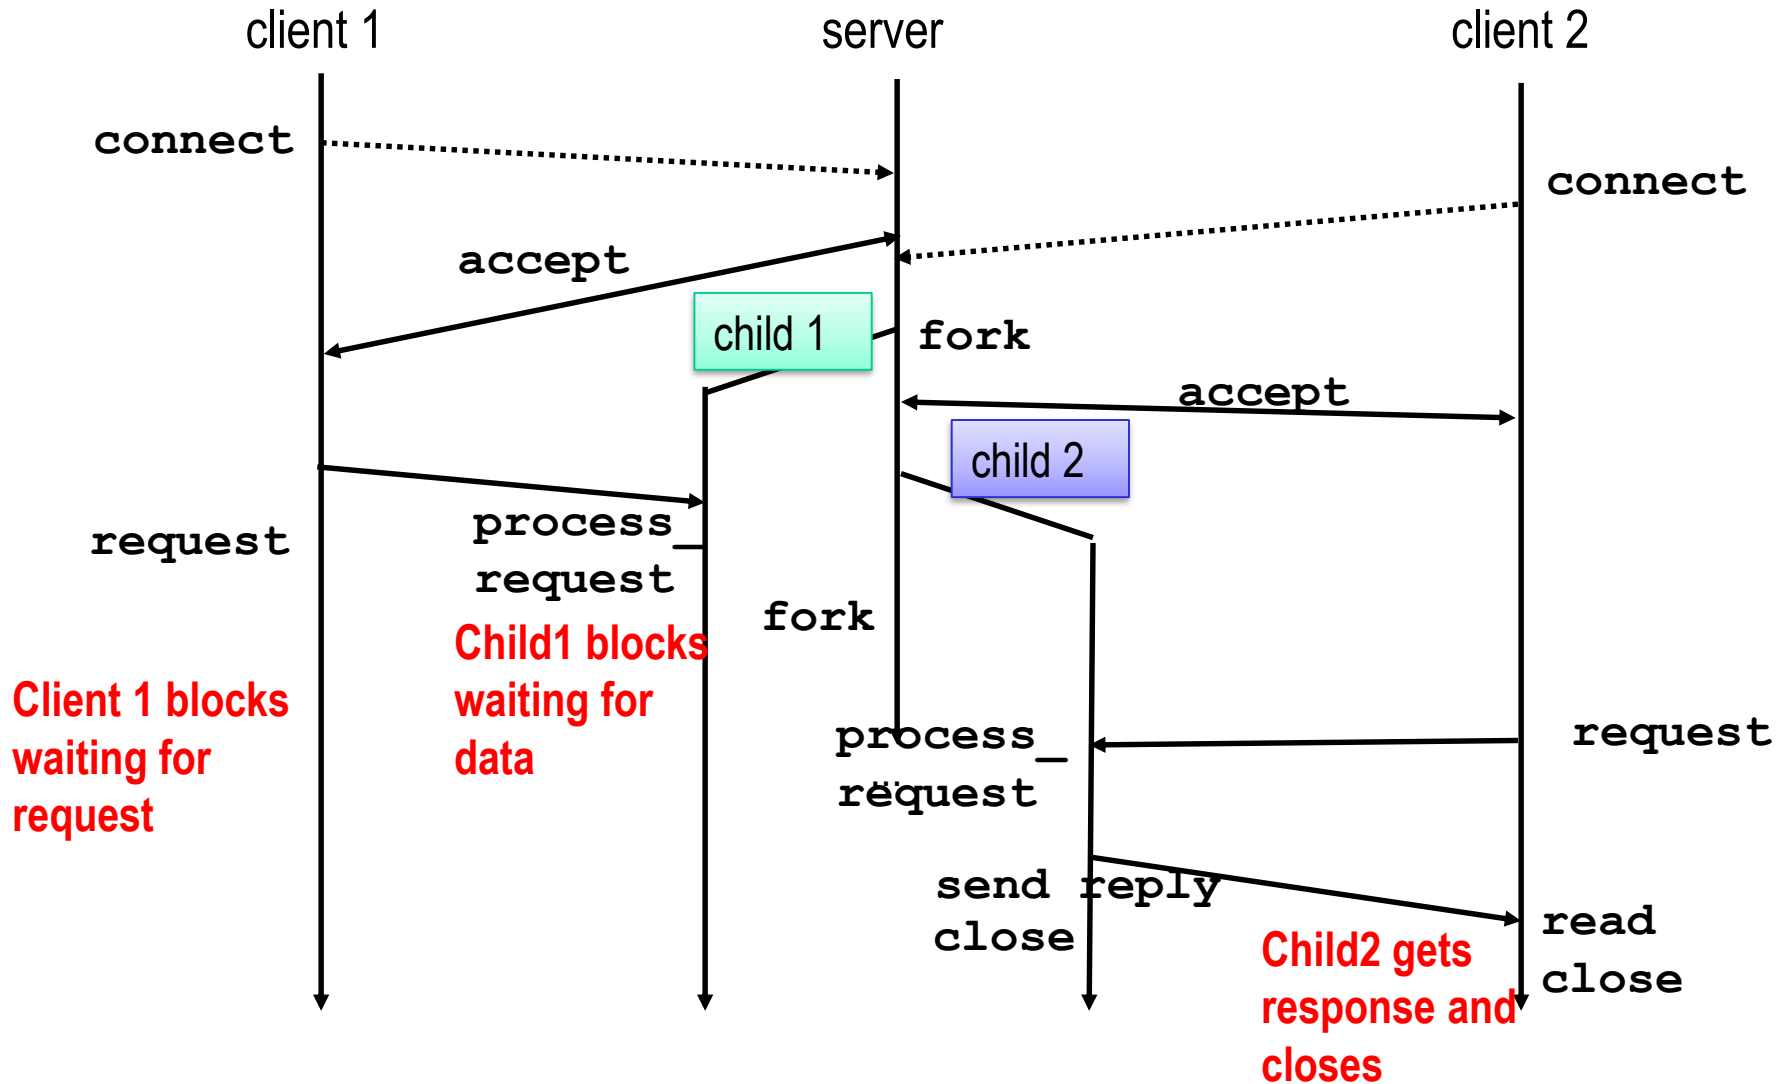
- Concurrent servers use multiple concurrent flows to serve multiple clients at the same time

The Simple Solution

```
int pid;
for (;;) {
    receive(&client, request);
    if ((pid = fork()) == 0) {
        process_request(...);
        send (client, reply);
        _exit(0); // done
    } // if
} // for
```


The Simple Solution (con't)


- Spawn separate process for each client



fork Example: Two consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye


Infeasible output:

L0
Bye
L1
Bye
L1
Bye
Bye

fork Example: Nested forks in parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
L1
Bye
Bye
L2
Bye


Infeasible output:

L0
Bye
L1
Bye
Bye
L2

fork Example: Nested forks in children

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
Bye
L1
L2
Bye
Bye

Infeasible output:

L0
Bye
L1
Bye
Bye
L2

Today

- System call error handling
- More on fork System Call and Process Graphs
- **wait and waitpid system calls**
- **exec family system calls**
- **signals**

Reaping Child Processes

■ Problem

- When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
- Called a “zombie”
 - Living corpse, half alive and half dead



■ Reaping

- Performed by parent on terminated child (using `wait` or `waitpid`)
- Parent is given exit status information
- Kernel then deletes zombie child process

■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
- So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n", getpid());  
        exit(0);  
    } else {  
        printf("Running Parent, PID = %d\n", getpid());  
        while (1); /* Infinite loop */  
    }  
}
```

forks.c

```
linux> ./forks 7 &  
[1] 6639  
Running Parent, PID = 6639  
Terminating Child, PID = 6640
```

```
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6639 ttyp9        00:00:03 forks  
 6640 ttyp9        00:00:00 forks <defunct>  
 6641 ttyp9        00:00:00 ps
```

```
linux> kill 6639  
[1]    Terminated
```

```
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6642 ttyp9        00:00:00 ps
```

■ **ps** shows child process as “defunct” (i.e., a zombie)

■ Killing parent allows child to be reaped by **init**

Zombie Example

- `ps -t` shows child process as “defunct”
- Killing parent allows child to be reaped

First letter:

S: sleeping

T: stopped

R: running/ runnable

Z: Zombie

```

UNIX% ./example &
[1] 11299
Running Parent, PID = 11299
Terminating Child, PID = 11300
UNIX% ps x
  PID TTY          STAT TIME COMMAND
 11263 pts/7        Ss   0:00 -tcsh
 11299 pts/7        R    0:07 ./example
 11300 pts/7        Z    0:00 [...] <defunct>
 11307 pts/7        R+   0:00 ps x
UNIX% kill 11299
[1] Terminated
UNIX% ps x
  PID TTY          STAT TIME COMMAND
 11263 pts/7        Ss   0:00 -tcsh
 11314 pts/7        R+   0:00 ps x

```

Second letter:

s: session leader

+: foreground proc group

Refer to man page for details

```

D    uninterruptible sleep (usually IO)
I    Idle kernel thread
R    running or runnable (on run queue)
S    interruptible sleep (waiting for an event)
T    stopped by job control signal
t    stopped by debugger during the tracing
W    paging (not valid since the 2.6.xx kernel)
X    dead (should never be seen)
Z    defunct ("zombie") process, terminated but
    its parent

```


Non-terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n", getpid());
        while (1); /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n", getpid());
        exit(0);
    }
}
```

forks.c

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6676 tttyp9      00:00:06 forks
 6677 tttyp9      00:00:00 ps
linux> kill -9 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6678 tttyp9      00:00:00 ps
```

■ Child process still active even though parent has terminated

■ Must kill child explicitly, or else will keep running indefinitely

`wait`: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the `pid` of the child process that terminated
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`
 - See textbook for details

`wait`: Synchronizing with Children

You can call them after calling `wait` to learn the reason of changing state


- `WIFEXITED(status)`: Returns true if the child terminated normally, via a call to `exit` or a return.
- `WEXITSTATUS(status)`: Returns the exit status of a normally terminated child. This status is only defined if `WIFEXITED` returned true.
- `WIFSIGNALED(status)`: Returns true if the child process terminated because of a signal that was not caught. (Signals are explained in Section 8.5.)
- `WTERMSIG(status)`: Returns the number of the signal that caused the child process to terminate. This status is only defined if `WIFSIGNALED(status)` returned true.
- `WIFSTOPPED(status)`: Returns true if the child that caused the return is currently stopped.
- `WSTOPSIG(status)`: Returns the number of the signal that caused the child to stop. This status is only defined if `WIFSTOPPED(status)` returned true.

wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

forks.c

Process graph is a useful tool for capturing the partial ordering of statements in a concurrent program:



Feasible output-1:

HC
HP
CT
Bye

Feasible output-2:

HP
HC
CT
Bye

Infeasible output:

HP
CT
Bye
HC

Another wait Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10() {  
    pid_t pid[N];  
    int i, child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0) {  
            exit(100+i); /* Child */  
        }  
  
    for (i = 0; i < N; i++) { /* Parent */  
        pid_t wpid = wait(&child_status);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

pid[0]=224

pid[1]=123

pid[2]=512

PID
224

→ exit(100)

PID
123

→ exit(101)

PID
512

→ exit(102)

Assume N = 3

WIFEXITED(status)

returns true if the child
terminated normally

WEXITSTATUS(status)

returns the exit status of the child.

Second for loop will run until all
children are reaped in **arbitrary
order**

Another wait Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10() {  
    pid_t pid[N];  
    int i, child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0) {  
            exit(100+i); /* Child */  
        }  
    for (i = 0; i < N; i++) { /* Parent */  
        pid_t wpid = wait(&child_status);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

pid[i]	PID	exit()	wpid
pid[0]=224	PID 224	exit(100)	wpid = 224
pid[1]=123	PID 123	exit(101)	wpid = 123
pid[2]=512	PID 512	exit(102)	wpid = 512

Child 224 is terminated with exit status 100
Child 512 is terminated with exit status 102
Child 123 is terminated with exit status 101

waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Various options available

```
void fork11() {  
    pid_t pid[N];  
    int i;  
    int child_status;  
  
    pid[0]=224  
    pid[1]=123  
    pid[2]=512  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0)  
            exit(100+i); /* Child */  
    for (i = N-1; i >= 0; i--) {  
        pid_t wpid = waitpid(pid[i], &child_status, 0);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

Assume N = 3

WIFEXITED(status)

returns true if the child
terminated normally

WEXITSTATUS(status)

returns the exit status of the child.

Second for loop will run until all
children are reaped in order from
last to first

waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Various options available

```
void fork11() {  
    pid_t pid[N];  
    int i;  
    int child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0)  
            exit(100+i); /* Child */  
    for (i = N-1; i >= 0; i--) {  
        pid_t wpid = waitpid(pid[i], &child_status, 0);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n",  
                wpid);  
    }  
}
```

The diagram illustrates the execution of the `waitpid` function in the `fork11` function. It shows three child processes, each identified by a PID (224, 123, and 512) and their corresponding exit status (100, 101, and 102). The `waitpid` function is called for each child, and the `wpid` variable is set to the PID of the terminated child. The `waitpid` function is highlighted in green in the code, and the `exit` calls are highlighted in red in the diagram.

Child 512 is terminated with exit status 102
Child 123 is terminated with exit status 101
Child 224 is terminated with exit status 100

waitpid options

The value of *options* is an OR of zero or more of the following constants:

WNOHANG

return immediately if no child has exited.

WUNTRACED

also return if a child has stopped

WCONTINUED (since Linux 2.6.10)

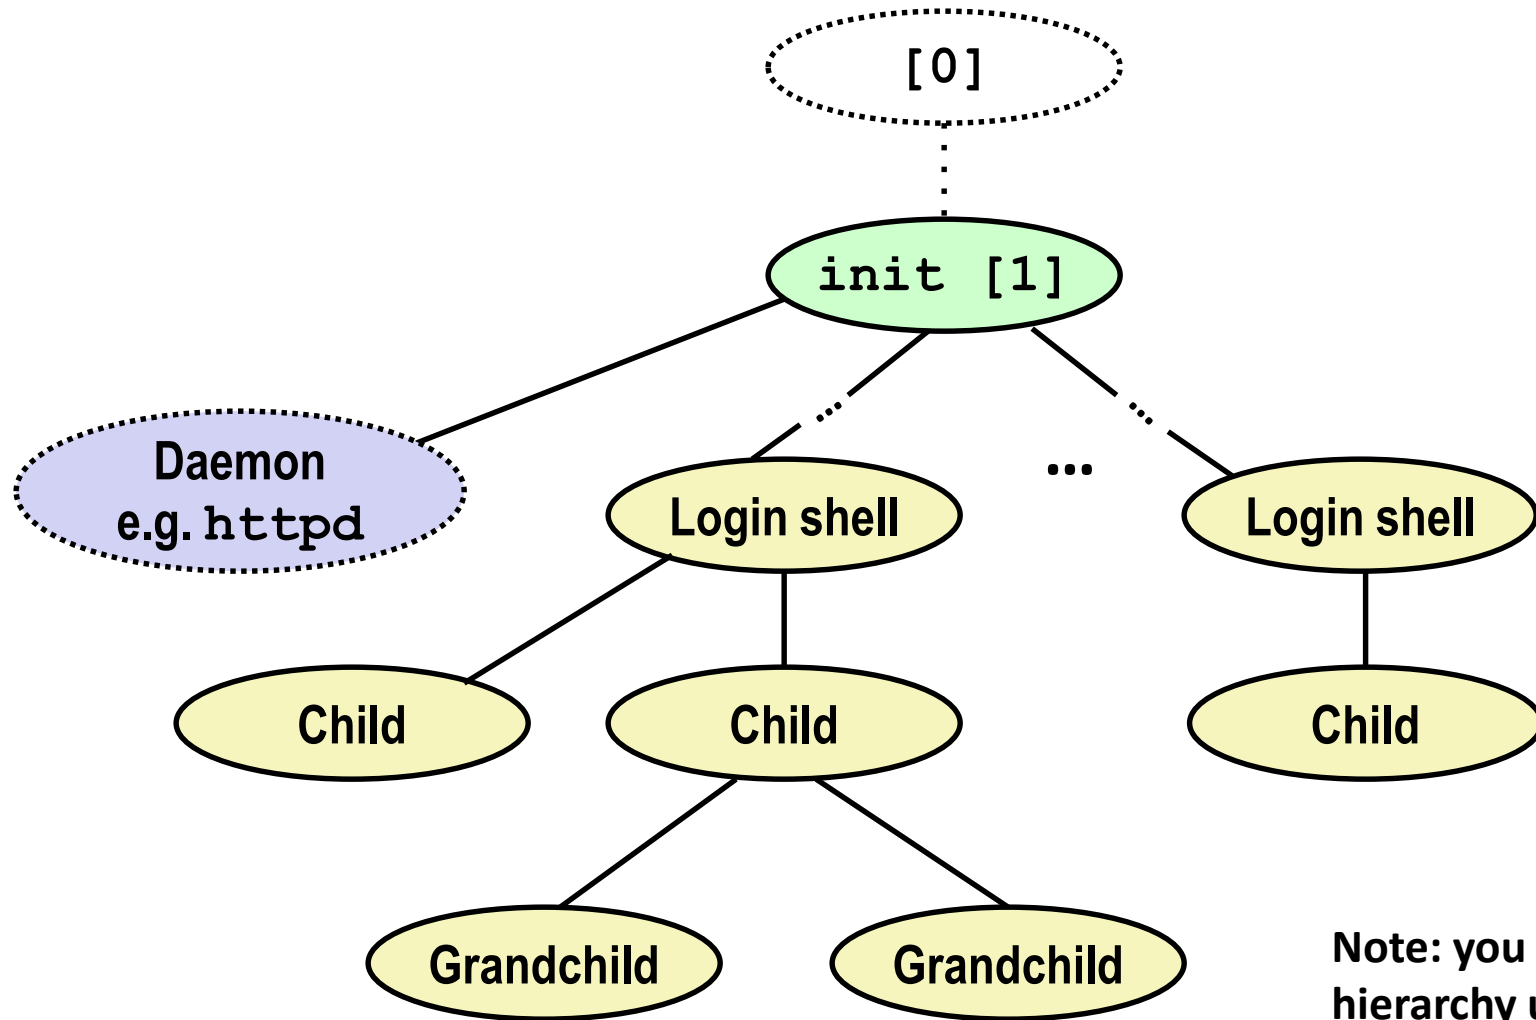
also return if a stopped child has been resumed by delivery of **SIGCONT**.
(For Linux-only options, see below.)

Use man for details

Today

- System call error handling
- More on fork System Call and Process Graphs
- wait and waitpid system calls
- **exec family system calls**
- **signals**


Linux Process Hierarchy



Note: you can view the hierarchy using the Linux `ps tree` command

UNIX Startup: 1


- Pushing reset button loads the PC with the address of a small bootstrap program
- Bootstrap program loads the boot block (disk block 0)
- Boot block program loads kernel from disk
- Boot block program passes control to kernel
- Kernel handcrafts the data structures for process 0



Process 0: handcrafted kernel process

Process 1: user mode process
`fork()` and `exec(/sbin/init)`


UNIX Startup: 2



`init` forks new processes as per the `/etc/inittab` file


Forks `getty` (get tty or get terminal) for the console

UNIX Startup: 3



`getty` execs a login program

UNIX Startup: 4



`login` gets user's uid & password

- If OK, it execs appropriate shell
- If not OK, it execs `getty`

pstree

```
systemd--NetworkManager--2*[{NetworkManager}]
--agetty
--auditd--{auditd}
--avahi-daemon--avahi-daemon
--chronyd
--crond
--dbus-daemon--{dbus-daemon}
--firewalld--{firewalld}
--irqbalance
--lvmetad
--master--pickup
--qmgr
--polkitd--6*[{polkitd}]
--python--2*[{python}]
--qemu-ga
--rsyslogd--2*[{rsyslogd}]
--smartd
--sshd--4*[sshd--sshd--bash--vim]
--sshd--sshd--sftp-server
```

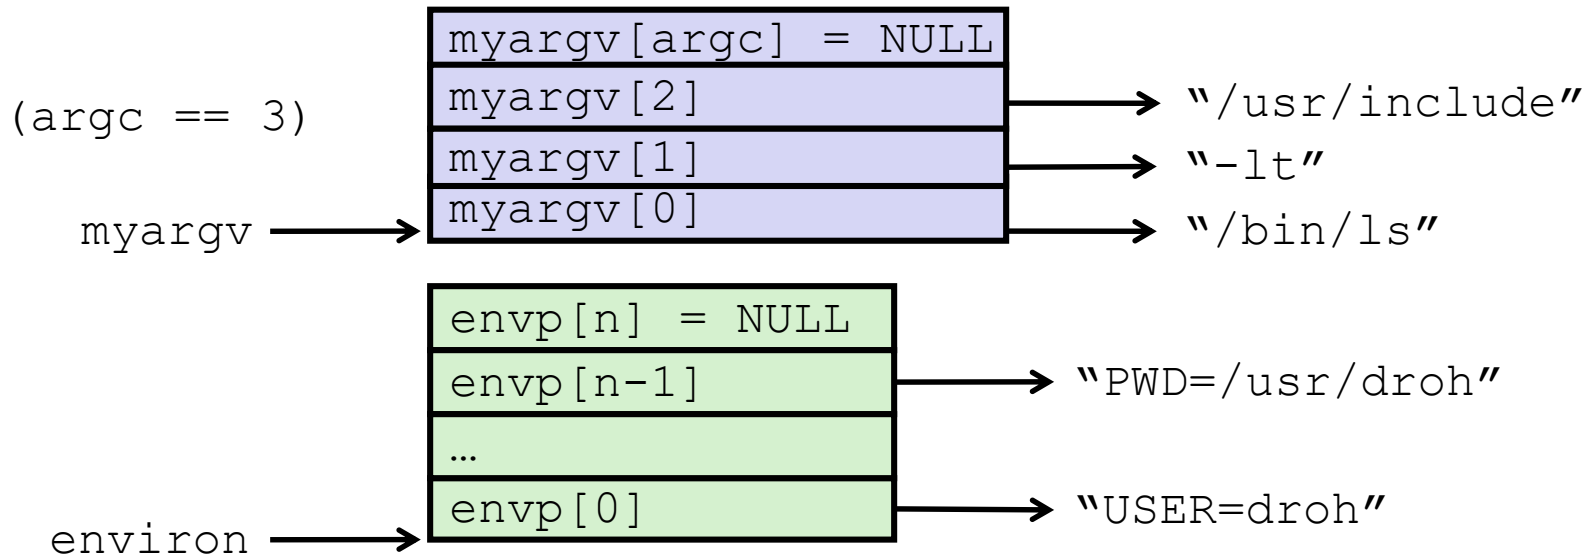
```
--sshd--4*[sshd--sshd--bash--vim]
--sshd--sshd--sftp-server
--sshd--sshd
--sshd--sshd--bash--pstree
--2*[sshd--sshd--bash]
```


execve : Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- **Loads and runs in the current process:**
 - Executable file `filename`
 - Can be object file or script file beginning with `#!interpreter` (e.g., `#!/bin/bash`)
 - ...with argument list `argv`
 - By convention `argv[0]==filename`
 - ...and environment variable list `envp`
 - “name=value” strings (e.g., `USER=droh`)
- **Overwrites code, data, and stack**
 - Retains PID, open files and signal context
- **Called **once** and **never** returns**
 - ...except if there is an error

execve Example

- Executes `"/bin/ls -lt /usr/include"` in child process using current environment:



```
if ((pid = Fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```

exec family function varieties

- **execl**

- **execle**

- **execlp**

- **execv**

- **execve**

- **execvp**

e(environment): It is an array of pointers that points to environment variables and is passed explicitly to the newly loaded process.

l: l is for the command line arguments passed a list to the function

p: p is the path environment variable which helps to find the file passed as an argument to be loaded into process.

v: v is for the command line arguments. These are passed as an array of pointers to the function.

Shell Programs

- A *shell* is an application program that runs programs on behalf of the user.

- `sh` Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
- `csch/tcsh` BSD Unix C shell
- `bash` “Bourne-Again” Shell (default Linux shell)

```
int main()
{
    char cmdline[MAXLINE]; /* command line */

    while (1) {
        /* read */
        printf("> ");
        fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
}
```

shellex.c

*Execution is a
sequence of
read/evaluate
steps*

Simple Shell eval Function

```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = Fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitig: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

**parseline,
builtin_command not
shown**

parseline updates argv
array and returns if process
is background

builtin_command, check if
one of the built in
commands: quit

child is running the user job

Only reap non-background
jobs

Problem with Simple Shell Example

- Our example shell correctly waits for and reaps foreground jobs
- But what about background jobs?
 - Will become zombies when they terminate
 - Will never be reaped because shell (typically) will not terminate
 - Will create a memory leak that could run the kernel out of memory

Running a job in background

- Add & to the end
- Or suspend a process running in foreground using CTRL+Z

```
[sonmeza@cmssc257 ~]$ ./background &
[1] 166632
[sonmeza@cmssc257 ~]$ ps
  PID TTY          TIME CMD
 66338 pts/10        00:00:00 bash
 66632 pts/10        00:00:05 background
 66637 pts/10        00:00:00 ps
[sonmeza@cmssc257 ~]$
```

- When you exit and start another session, you will see background is still running.

```
166632 sonmeza 20 0 4208 352 276 R 100.0 0.0 12:16.32 background
167007 admin 20 0 201748 66056 5200 D 10.0 0.1 0:00.50 perl.pl
```

ECF to the Rescue!

■ Solution: Exceptional control flow

- The kernel will interrupt regular processing to alert us when a background process completes
- In Unix, the alert mechanism is called a *signal*

Programmer's Model of Multitasking

■ Basic Functions

- `fork()` spawns new process
 - Called once, returns twice
- `exit()` terminates own process
 - Called once, never returns
 - Puts process into “zombie” status
- `wait()` and `waitpid()` wait for and reap terminated children
- `exec1()` and `execve()` run a new program in an existing process
 - Called once, (normally) never returns

■ Programming Challenge

- Understanding the nonstandard semantics of the functions
- Avoiding improper use of system resources
 - E.g., “Fork bombs” can disable a system

Today

- System call error handling
- More on fork System Call and Process Graphs
- wait and waitpid system calls
- exec family system calls
- **Signals**
 - Will be covered next